
Arab Republic of Egypt
Ministry of Communications and Information Technology
Digital Pioneers Initiative (Digilians)

GRADUATION PROJECT REPORT
**Professional Diploma in Applied Artificial Intelligence
and Data Analytics**

End-to-End Transportation & Logistics Tracking Intelligence System

*A Dual-System Architecture: Warehouse Intelligence (Power BI)
& Predictive Delivery Analytics (Python | SQL | Orange)*

Prepared By

Aml Wajeh Mamdouh Abuzeid | Mai Mohamed Ibrahim Beltagy | Manar Harby Abd Elmoneam Ali |
Alaa ahmed AbdelBasset Elgendy | Gannatallah Emad Abdelkader Elsayed

Under the Supervision of

DR/ Mohamed Sobhy Mahmoud Mohamed

March 2026

Digital Pioneers Initiative (Digilians)— Applied AI & Data Analytics Program

ACKNOWLEDGMENTS

On behalf of the entire project team, we wish to express our profound and sincere gratitude to all those whose guidance, support, and institutional commitment made the completion of this work possible.

The Ministry of Communications and Information Technology — Digital Pioneers Initiative (Digilians)

We extend our deepest appreciation to the Ministry of Communications and Information Technology of the Arab Republic of Egypt for conceiving and delivering the Digital Pioneers Initiative (Digilians)— a programme of genuine national significance. By investing in the analytical and technical capabilities of Egypt's emerging generation of data professionals, the Ministry has created a learning environment that is both intellectually rigorous and practically transformative. The structured curriculum, industry-aligned faculty, and institutional infrastructure provided through this programme were indispensable to the depth and quality of the analytical work documented in this report. We are honoured to have been participants in this initiative and deeply conscious of the responsibility it carries.

Project Supervisor

We owe a particular debt of gratitude to our project supervisor, DR/ Mohamed Sobhy Mahmoud Mohamed whose expertise, intellectual generosity, and sustained engagement were central to the development of this project. Their guidance in framing the analytical problem, validating our methodological choices, and challenging us to achieve greater precision and clarity in our outputs elevated the quality of this work in ways that would not have been possible without their involvement. We are sincerely thankful for the time, insight, and professional mentorship they so willingly invested throughout the lifecycle of this project.

The Project Team

This project was the product of sustained collective effort, intellectual collaboration, and shared commitment to analytical excellence. Each member of the team contributed with diligence, creativity, and a genuine dedication to delivering a system worthy of the standards set by this programme. The challenges encountered — in database engineering, machine learning pipeline design, visual analytics, and report synthesis — were met with rigour and perseverance. We are proud of what we have built together, and deeply grateful for the trust, effort, and professionalism that every team member brought to this undertaking.

Finally, we wish to acknowledge the contributions of the broader academic and professional community whose published research, open-source tools, and shared knowledge informed the technical foundations of this system. The fields of applied data science, supply chain analytics, and machine learning are built on the collective intellectual generosity of many — and this project is one more addition to that ongoing endeavour.

EXECUTIVE SUMMARY

Business Problem | The Challenge of Reactive Logistics Management

The logistics and supply chain operation that forms the subject of this project exhibited a set of deeply interconnected operational deficiencies that collectively defined a network operating below its performance potential and above its risk tolerance. Chief among these was the absence of any predictive risk management capability: under the prevailing operational model, a shipment was classified as delayed only after its contractual Service Level Agreement had been breached — at which point the window for effective intervention had already closed. Compounding this structural reactive posture were critical inventory visibility gaps in the inbound warehouse environment, a financially material Order Fulfillment Rate deficit of 10.03 percentage points against the 95% industry benchmark, systemic vendor non-performance concentrated in specific high-volume accounts, and an Estimated Time of Arrival calculation framework demonstrably inadequate for long-haul freight corridors exceeding 1,000 kilometres. The cumulative financial exposure attributable to these deficiencies was estimated at approximately \$38.67 million per month in unrecovered revenue — a figure that made the development of a comprehensive intelligence response not merely a strategic opportunity, but an operational imperative.

The Solution | A Dual-System End-to-End Intelligence Architecture

In response to these challenges, this project designed, developed, and validated an End-to-End Transportation and Logistics Tracking Intelligence System built on a Dual-System analytical architecture. The first system — the Inbound Analytics Layer — was constructed as a Power BI enterprise intelligence dashboard analysing 3,204 SKUs across five product categories and four warehouse zones, delivering quantified financial impact assessment, inventory risk profiling, and zone-level operational benchmarking. The second system — the Outbound Analytics Layer — was powered by a Microsoft SQL Server relational database (the LogisticT schema: eleven normalised tables and ten Business Intelligence SQL VIEWS), a Python-driven predictive pipeline incorporating Haversine-based geospatial imputation, Exploratory Data Analysis, Apriori association rule mining, and a supervised Random Forest Classifier trained to assign delay risk probabilities to every new shipment at the moment of booking. Supplementary visual data mining conducted in Orange Data Mining surfaced fleet performance diagnostics, temporal demand characterisation, and ETA feasibility analysis. The entire architecture rests on a constraint-enforced relational foundation designed to Third Normal Form — ensuring that every dashboard visualisation and every machine learning prediction is grounded in structurally and semantically guaranteed data quality.

Value Delivered | From Fragmented Data to Actionable Predictive Intelligence

The primary value delivered by this system is the transformation of a fragmented, siloed, and underutilised operational data estate into a unified, 360-degree intelligence platform that serves the warehouse manager, the dispatch supervisor, the financial controller, and the executive leadership team with equal analytical rigour. The organisation gains, for the first time, the capability to anticipate delivery risk before it materialises, allocate fleet and warehouse resources with statistical precision, hold vendors accountable against evidence-backed performance benchmarks, and engage customers with the proactive transparency that converts service reliability into sustainable commercial loyalty. This is not an incremental improvement to existing reporting capabilities —

it is a foundational transformation in how the organisation understands, monitors, and governs its supply chain: establishing the analytical infrastructure required to lead, rather than follow, in an industry where intelligence-led operations are rapidly becoming the competitive baseline.

KEY SYSTEM METRICS AT A GLANCE			
3,204 SKUs analysed		237 Critical SKUs below reorder point	~\$38.67M monthly revenue at risk
11-table SQL schema		10 BI SQL VIEWS	Random Forest Classifier (5-fold stratified cross-validation)
Technology Stack: Microsoft SQL Server Power BI Python (pandas, scikit-learn) Orange Data Mining			

CHAPTER 1 : Introduction

The global logistics and supply chain industry stands today at the intersection of unprecedented complexity and transformative technological opportunity. As international trade volumes continue to expand and consumer expectations for delivery speed and transparency reach new heights, the organizations that manage the physical movement of goods face an existential imperative: **evolve from reactive, experience-driven operations into data-driven, intelligence-led systems**, or risk being outpaced by competitors who do. This report documents the design, development, and deployment of a comprehensive analytics and artificial intelligence system purpose-built to address that imperative.

The project, titled *End-to-End Transportation & Logistics Tracking Intelligence System*, was developed as the capstone deliverable for the **Professional Diploma in Applied Artificial Intelligence and Data Analysis** — delivered through the **Digital Egypt Vanguard**s initiative. It represents a complete, enterprise-grade solution that integrates relational database engineering, business intelligence dashboarding, machine learning, and visual data mining into a single unified analytical platform serving both the inbound warehouse environment and the outbound last-mile delivery network.

1.1 The Modern Logistics Industry: Scale, Complexity & the Data Imperative

1.1.1 The Global Context

Logistics and supply chain management constitutes one of the most economically significant and operationally complex sectors of the modern global economy. The physical movement of goods — from raw material procurement through manufacturing, warehousing, and final-mile customer delivery — underpins virtually every commercial transaction in the world. At scale, logistics networks encompass hundreds of thousands of daily shipment events, multi-tier vendor hierarchies, real-time fleet deployments, and dynamic customer demand patterns that shift in response to market conditions, seasonal cycles, and macroeconomic forces.

The operational stakes within this environment are formidable. A single SLA breach — a shipment delivered late against its contractual commitment — triggers a cascade of consequences: direct financial penalties, customer dissatisfaction, potential loss of commercial contracts, and the reputational damage that accompanies repeated service failures. Conversely, organizations that achieve consistent, transparent, and predictable delivery

performance command premium customer loyalty, higher contract renewal rates, and a sustainable competitive differentiation that is increasingly difficult to replicate without an equivalent data infrastructure.

1.1.2 The Technology Shift: From Spreadsheets to Intelligence Systems

The logistics industry has historically relied on experience-based decision-making, manual exception handling, and retrospective performance reporting — tools adequate for simpler, lower-volume operating environments but demonstrably insufficient for the scale and velocity of modern supply chains. The emergence of cloud-scale databases, scalable business intelligence platforms, advanced machine learning frameworks, and GPS telematics has created both the necessity and the technical means for a fundamental shift in how logistics operations are monitored, managed, and optimized.

The core transformation this project addresses can be summarized in a single strategic phrase: the shift from **reactive logistics management** — in which problems are identified only after they occur — to **proactive, intelligence-led operations** — in which risks are quantified in advance, patterns are detected before they become failures, and every operational decision is grounded in analytical evidence rather than intuition. This transformation requires not merely the adoption of individual analytical tools, but the architectural integration of data engineering, visualization, and predictive modeling into a coherent, end-to-end intelligence platform.

1.2 Project Scope: The Dual-System Architecture

A defining characteristic of this project — and a primary contributor to its analytical ambition — is its **dual-system architectural approach**. Rather than addressing either the inbound warehouse environment or the outbound delivery network in isolation, the system was designed to operate across both operational dimensions simultaneously, recognizing that effective last-mile delivery performance is inseparable from the warehouse readiness that precedes it.

1.2.1 System 1: Inbound Analytics — Warehouse Intelligence (Power BI)

The first system of the dual architecture is directed inward — at the warehouse and inventory management environment that constitutes the logistical origin point for every outbound delivery. Built on a **Power BI enterprise intelligence dashboard** analyzing a dataset of **3,204 Stock-Keeping Units (SKUs)** across five product categories and four warehouse zones, this system delivers three categories of analytical intelligence:

- **Financial Impact Assessment:** Quantifying the revenue exposure created by the current Order Fulfillment Rate of **84.97%** — a 10.03 percentage-point deficit against the 95% industry benchmark, representing an estimated **~\$38.67 million per month** in revenue at risk.
- **Inventory Risk Profiling:** Identifying and classifying the **237 SKUs (7.4% of the total portfolio)** currently operating below their designated reorder point, and generating category-level stockout risk profiles for the Automotive and Pharma segments.
- **Operational Zone Benchmarking:** Benchmarking picking time and throughput performance across all four warehouse zones against Zone A as the internal performance standard, identifying specific operational inefficiencies and their measurable impact on dispatch scheduling.

1.2.2 System 2: Outbound Analytics — Last-Mile Delivery Tracking (Python & SQL)

The second system is directed outward — at the road network, carrier performance, and last-mile delivery execution that ultimately determines the customer experience. This system is powered by a **Python-based analytical pipeline** and a **Microsoft SQL Server relational database** (the **LogisticT** schema) comprising eleven normalized tables and a suite of ten Business Intelligence SQL VIEWS. Its analytical outputs span four domains:

- **Predictive Delay Classification:** A supervised **Random Forest Classifier** trained to assign a delay risk probability to every new shipment at the moment of booking — before the vehicle departs — enabling pre-emptive operational interventions that prevent SLA breaches rather than managing them after the fact.
- **Geospatial Route Analytics:** Application of the **Haversine Formula** for GPS-based distance computation, and K-Means clustering to identify route performance archetypes and anomalous trip profiles.
- **SLA Breach & Courier Performance Monitoring:** SQL-driven views providing precise delay-hour quantification, courier-level success rate scoring, and real-time identification of shipments with no GPS movement for over 48 hours.
- **Visual Data Mining (Orange):** A four-component visual analytics investigation into fleet performance distributions, daily demand rhythms, vendor delay patterns, and the correlation between route distance and delay probability — culminating in a finding that the current ETA model is structurally inadequate for routes exceeding 1,000 km.

1.2.3 The Unifying Logic: Why Both Systems Are Necessary

The dual-system architecture is not a design choice born of analytical ambition alone — it reflects a fundamental operational truth that the data in this project consistently validates: **warehouse performance and delivery performance are not independent variables**. The stockout events, elevated picking times, and fulfillment gaps identified in the inbound Power BI system are direct upstream causes of the dispatch delays, route volatility, and ML model accuracy limitations that manifest in the outbound delivery network.

As demonstrated in Chapter 4 of this report, a warehouse operating with high stockout rates generates irregular dispatch patterns that degrade the historical data quality on which the outbound predictive models depend. Conversely, a warehouse achieving Zone A-level performance standards produces the stable, consistent dispatch rhythms that enable high-confidence machine learning ETA predictions. **Warehouse optimization is therefore a precondition for delivery intelligence — and this project's dual architecture makes that causal relationship explicit, measurable, and actionable.**

1.3 Technology Stack & Methodological Framework

The system was built using an integrated stack of industry-standard data engineering and analytics technologies, each selected for its alignment with the specific analytical requirements of its respective system component:

Technology	System Layer	Primary Analytical Role
Microsoft SQL Server	Data Engineering	Relational database schema design, normalization, constraint enforcement, and 10 BI VIEWS for the LogisticT system
Microsoft Power BI	Inbound Analytics	Interactive KPI dashboards for inventory risk, zone benchmarking, and fulfillment gap analysis across 3,204 SKUs
Python (pandas, scikit-learn, matplotlib, seaborn)	Outbound Analytics & ML	Data engineering, EDA, K-Means clustering, association rule mining, and Random Forest delay classification
Orange Data Mining	Visual Analytics	No-code visual data mining for fleet performance analysis, vendor SLA profiling, and ETA model validity assessment
Haversine Formula	Geospatial Engineering	GPS coordinate-based distance computation for missing data imputation and route distance accuracy assurance
Random Forest Classifier	Predictive Modeling	Supervised delay risk scoring at booking time using 13 engineered features; validated via 5-fold stratified cross-validation

Table 1.1 — Integrated Technology Stack: End-to-End Transportation & Logistics Tracking Intelligence System

1.4 Report Structure

This report is organized into seven chapters, each addressing a distinct phase of the project's analytical lifecycle. The chapters are designed to be read sequentially, as each builds upon the methodological and analytical outputs of its predecessors:

Chapter	Title	Content Summary
Chapter 1	Introduction	Context, scope, dual-system architecture overview, technology stack, and report structure
Chapter 2	Business Problem & Objectives	Operational pain points, strategic objectives, and the analytical transformation goals of the project
Chapter 3	Database Architecture & Business Intelligence (SQL)	LogisticT SQL Server schema design, normalization, constraint engineering, and 10 BI SQL VIEWS
Chapter 4	Inbound Analytics & Warehouse Optimization (Power BI)	Inventory risk profiling, fulfillment gap analysis, zone benchmarking, and the inbound-to-outbound causal link
Chapter 5	Predictive Analytics & Data Mining (Python)	Feature engineering, EDA, K-Means clustering, association rule mining, and Random Forest delay classification
Chapter 6	Visual Data Mining (Orange)	Prioritised, evidence-based operational improvement recommendations addressed to executive leadership
Chapter 7	Strategic Business Recommendations	Fleet performance analysis, temporal demand patterns, vendor SLA profiling, and ETA model recalibration findings

Table 1.2 — Report Chapter Map

Chapter Summary

This chapter has established the strategic and contextual foundation for the project. The modern logistics industry demands an analytical infrastructure capable of converting operational data into real-time intelligence — and this project delivers that infrastructure through a dual-system architecture that addresses both the inbound warehouse environment and the outbound delivery network. The chapters that follow document how that architecture was designed, built, and validated against real operational data.

CHAPTER 2 :Business Problem & Objectives

Every analytical system is an answer to a question. Before documenting the technical architecture and analytical outputs of this project, it is essential to precisely define **the operational problems it was designed to solve** — and the strategic objectives it was designed to achieve. This chapter articulates the specific business pain points that motivated the development of the dual-system intelligence platform, and frames the project's goals in terms of measurable, outcome-oriented objectives that are directly addressable through data analytics and artificial intelligence.

The problems documented here were identified through analysis of the operational dataset underpinning this project. They are not generic industry challenges — they are **quantified, evidence-backed deficiencies** observed in the specific logistics network under investigation, with financial and operational impacts that are calculable, attributable, and ultimately preventable through the intelligent application of the methodologies deployed in subsequent chapters.

2.1 The Business Problem: A Network Operating Without Intelligence

The logistics operation analyzed in this project exhibited a consistent and interrelated set of operational deficiencies that, taken collectively, define a network **operating below its performance potential and above its risk tolerance**. Four primary problem domains were identified:

2.1.1 Problem 1: Reactive Logistics — The Absence of Predictive Risk Management

The most strategically consequential operational deficiency identified in this project is structural and cultural rather than technical: the logistics network operates on a **fundamentally reactive model**. Under the current operational paradigm, a shipment is classified as delayed only *after* its committed delivery SLA has been breached — at which point the operational window for effective intervention has closed. Customer communication, re-routing decisions, and compensatory measures must be initiated under time pressure, with limited options, and at significantly higher cost than pre-departure intervention would have required.

This reactive posture is not a product of analytical incapacity. It is a product of the **absence of a predictive risk scoring system** that can evaluate the delay probability of a shipment *before the vehicle leaves the warehouse*. The data required to make such predictions exists within the operational database. The machine learning

methodology capable of extracting predictive signal from that data is well-established. What is absent is the integrated analytical infrastructure that connects them — the precise gap that this project's outbound analytics system is designed to fill.

Quantified Impact — Reactive Operations

Analysis of the outbound dataset revealed that the majority of delayed shipments share identifiable characteristics — route distance band, vehicle category, origin state, and booking time period — that are **predictable at the time of dispatch planning**. The Random Forest Classifier developed in Chapter 5 demonstrates that this predictive signal is statistically significant and practically exploitable, enabling the transformation from reactive failure management to proactive risk prevention.

2.1.2 Problem 2: SLA Breaches & Delivery Delay — Systemic and Structurally Caused

The outbound delivery dataset reveals a **systemic pattern of SLA breach** that is neither random nor uniformly distributed across the network. Late deliveries — defined as shipments where the actual delivery timestamp exceeds the contractually committed `PromisedDate` — are concentrated along specific route corridors, within specific vehicle categories, and against specific vendor accounts. This concentration is the analytical signature of **structural causation rather than stochastic operational variance**

Three structural causes were identified and quantified through the analytical pipeline documented in this report:

- **ETA Model Inadequacy for Long-Haul Routes:** Visual data mining analysis (Chapter 7) confirmed a positive and statistically significant correlation between route distance and delay probability. Shipments on routes exceeding 1,000 km exhibit markedly higher and more consistent delay rates than short-haul equivalents — not because carriers are underperforming, but because the ETA benchmarks against which they are measured were **never calibrated to reflect the operational reality of extended-distance freight movements**. Mandatory rest stop requirements, traffic variability, fuel and maintenance stops, and route-specific infrastructure constraints are systematically absent from current ETA calculations.
- **Vendor-Level Performance Failures:** Two vendor accounts — Otis Elevator (34 shipments, **100% delayed**) and Wipro (63 shipments, **98.4% delayed**) — exhibited delay rates that are categorically incompatible with SLA compliance. These figures indicate not operational variance but **contractual non-performance** that demands formal escalation through Performance Improvement Plans and, if unresolved, contract review.
- **Fleet Performance Differentials:** Heavy trailer vehicle categories exhibited systematically longer transit times on long-haul routes compared to lighter vehicle classifications, indicating fleet assignment policies that are not calibrated to route distance and time-sensitivity requirements.

Quantified Impact — SLA Breach Exposure

The `vw_LateDeliveryReport` SQL VIEW developed in Chapter 3 computes the precise **DelayHours** for every breached shipment — the difference between the `PromisedDate` and the actual `DeliveryDate`. This metric is the primary target variable for the Chapter 5 machine learning models, establishing a direct architectural connection between SLA breach quantification and predictive risk prevention. The current absence of a system capable of surfacing this risk before departure means every preventable delay currently represents a realized cost.

2.1.3 Problem 3: Fleet & Operational Visibility Gaps

The logistics operation currently lacks a **systematic and real-time framework** for monitoring fleet performance, courier effectiveness, and route efficiency. Several critical operational intelligence gaps were identified:

- **No Real-Time Stuck Shipment Detection:** Shipments that enter an in-transit or pending state and subsequently cease generating GPS movement logs — indicators of vehicle breakdowns, route blockages, or telematics failures — are not currently detected until a customer complaint or manual check occurs. In the dataset analyzed, **TRK_STUCK_001** had been stationary for over 48 hours at Vellore Toll with no operational alert triggered. The **vw_StuckShipments** view developed in this project addresses this gap with a live, self-updating 48-hour movement threshold.
- **No Courier Performance Scorecard:** Delivery agent performance — measured as the ratio of on-time and early deliveries against total assignments — was not systematically tracked or reported. Without this data, the organization cannot identify high-performing couriers for incentive programs, underperforming agents for targeted training, or workload distribution imbalances that may be contributing to delay patterns.
- **No Hub-Level Financial P&L Visibility:** The financial contribution of individual origin hubs — measured as shipping fee revenue net of vehicle maintenance costs — was not calculated or reported on a regular basis. This absence of hub-level profitability data prevents informed capital allocation decisions regarding hub investment, capacity adjustment, and route prioritization.
- **Incomplete Proof of Delivery (POD) Records:** GPS provider coverage for delivery confirmation events was inconsistent across the network, creating compliance gaps in the chain-of-custody documentation required for contractual and regulatory purposes.

2.1.4 Problem 4: Warehouse Inefficiencies — The Hidden Upstream Driver

Perhaps the most analytically significant finding of this project is the demonstration that inbound warehouse inefficiencies are not merely internal operational issues — they are **upstream causal drivers of outbound delivery failures**. The Power BI warehouse analytics system (Chapter 4) identified four warehouse-level conditions with direct, measurable impacts on last-mile delivery performance:

Inbound Problem	Quantified Magnitude	Outbound Consequence
Below-benchmark Order Fulfillment Rate	84.97% vs. 95% industry standard	Volatile dispatch patterns degrade ML model prediction accuracy
Critical SKUs below reorder point	237 SKUs (7.4% of portfolio)	Order fragmentation, holds, and increased re-dispatch events
High Automotive stockout frequency	3,110 stockout events recorded	Emergency procurement lag and unpredictable carrier scheduling
Zone picking time inefficiencies (B/C/D)	Above Zone A benchmark across three zones	Late truck dispatch windows compress last-mile delivery buffer

Table 2.1 — Inbound-to-Outbound Problem Causation Map

The inbound analytics system developed in this project transforms these previously invisible causal relationships into **visible, quantified, and actionable intelligence** — enabling warehouse managers and supply chain executives to understand not just that stockouts are occurring, but precisely how each stockout event degrades the reliability of the entire downstream delivery network.

2.2 Project Objectives

The four problem domains articulated in Section 2.1 define the analytical agenda of this project. The following objectives were established as the primary success criteria for the system design and implementation:

2.2.1 Objective 1 — Transform from Reactive to Proactive Decision-Making

The primary strategic objective of the outbound analytics system is to fundamentally re-architect the operational decision model from one of reactive failure management to one of proactive risk prevention. This objective is operationalized through the development and deployment of a **supervised machine learning delay risk classifier** that assigns a delay probability score to every new shipment at the moment of booking — prior to vehicle dispatch — enabling supervisors to trigger pre-emptive interventions including vehicle reassignment, route replanning, pre-emptive customer notification, and adjusted SLA commitment, all before a single kilometer has been driven.

Success against this objective is defined as the delivery of a trained, validated, and serializable Random Forest Pipeline with meaningful **Precision-Recall AUC performance** across both training and held-out test partitions, confirmed through 5-fold stratified cross-validation.

2.2.2 Objective 2 — Optimize ETA Prediction Accuracy & Recalibrate SLA Commitments

The analytical evidence from both the Python data mining pipeline (Chapter 5) and the Orange visual analytics system (Chapter 7) establishes that the current ETA calculation model is structurally flawed for long-haul routes. A core project objective is to **quantify the relationship between route distance and delay probability** through statistical and visual analysis, and to translate that quantification into actionable ETA model recalibration recommendations for routes exceeding defined distance bands (500–1,000 km, 1,000–1,500 km, and >1,500 km).

This objective also encompasses the identification and escalation of the **short-distance, high-duration outlier** cluster identified in Chapter 7 — shipments covering routes under 500 km but recording transit durations exceeding 250 hours — which represent operational efficiency failures distinct from the ETA model inadequacy and requiring personnel accountability responses rather than algorithmic adjustments.

2.2.3 Objective 3 — Minimize Operational Costs & Maximize Network Profitability

The project's financial objectives span both the inbound and outbound systems. On the inbound side, the objective is to **quantify and reduce the revenue exposure** created by the current 84.97% Order Fulfillment Rate, identify the specific SKU cohorts and product categories driving that exposure, and recommend targeted procurement and replenishment interventions that close the gap toward the 95% industry benchmark — a change that the Power BI analysis estimates would unlock access to a significant portion of the **\$257.80 million monthly revenue potential** currently inaccessible due to structural inventory inefficiencies.

On the outbound side, the objective is to surface **hub-level profitability intelligence** through the **vw_HubRevenueMonthly** SQL VIEW, enabling executives to compare shipping fee revenue against vehicle maintenance costs on a per-hub, per-month basis — and to identify underperforming or loss-generating hub-corridor combinations that warrant operational review or capital reallocation.

2.2.4 Objective 4 — Improve Customer Satisfaction Through Transparency & Reliability

Customer satisfaction in logistics is determined primarily by two factors: **delivery reliability** (arriving on time, as committed) and **delivery transparency** (knowing where a shipment is, and when it will arrive). The project addresses both dimensions:

- **Reliability:** The predictive delay classifier and ETA recalibration framework directly reduce the frequency of SLA breaches, while the vendor Performance Improvement Plan framework (Chapter 6) addresses the systemic under-performance of accounts exhibiting near-100% delay rates.
- **Transparency:** The **vw_TrackingTimeline** and **vw_LatestShipmentStatus** SQL VIEWS provide a real-time, customer-accessible shipment position feed. The **vw_PODValidation** VIEW ensures that GPS-confirmed proof of delivery records are consistently captured and available for compliance and dispute resolution purposes.
- **Proactive Communication:** By identifying high-risk shipments at booking time, the system enables the operations team to **proactively notify customers of elevated delay risk** before departure — transforming a reactive apology into a proactive service commitment that protects the customer relationship.

2.2.5 Objective 5 — Establish a Unified, Governed Data Infrastructure

Underpinning all four operational objectives above is a foundational technical objective: the design and deployment of a **constraint-enforced, normalized relational database** that serves as the single source of truth for all operational, financial, and geographical data in the logistics network. Without this data infrastructure — with its CHECK constraints, FOREIGN KEY referential integrity, DEFAULT values, and semantic Business Intelligence VIEW layer — the analytical systems built upon it would produce unreliable outputs regardless of their algorithmic sophistication.

This objective reflects a principle that the technical architecture of this project consistently validates: **analytical intelligence is only as trustworthy as the data governance that underlies it**. Data quality is not a pre-analytical concern to be handled upstream and forgotten — it is an ongoing architectural commitment that must be embedded into the structure of the database itself.

2.3 Summary: From Pain Points to Analytical Solutions

The relationship between the four problem domains identified in Section 2.1 and the five project objectives defined in Section 2.2 can be expressed as a structured mapping, which serves as the conceptual backbone of the entire analytical system:

Business Problem (Section 2.1)	Project Objective (Section 2.2)	System / Chapter Delivering the Solution
Reactive logistics; no predictive risk management	Obj. 1 — Proactive decision-making	Random Forest Classifier (Chapter 5); Dispatch Risk API (Chapter 6)
SLA breaches; flawed ETA model; vendor non-performance	Obj. 2 — ETA optimization & SLA recalibration	Visual Data Mining (Chapter 7); Python EDA (Chapter 5); SQL VIEWS (Chapter 3)
No hub P&L visibility; high fulfillment revenue gap	Obj. 3 — Cost minimization & profitability	Power BI Inbound Analytics (Chapter 4); vw_HubRevenueMonthly (Chapter 3)
Delivery unreliability; lack of tracking transparency	Obj. 4 — Customer satisfaction & transparency	SQL Tracking VIEWS (Chapter 3); POD Validation (Chapter 3)
No governed, trustworthy data foundation	Obj. 5 — Unified, governed data infrastructure	LogisticT SQL Server Schema (Chapter 3)

Table 2.2 — Problem-to-Objective-to-Solution Traceability Matrix

Chapter Conclusion & Bridge to Chapter 3

This chapter has defined the four business problem domains that motivated this project and the five strategic objectives that guided its design. Every technical decision documented in the chapters that follow — from the schema normalization strategy of the LogisticT database to the feature engineering choices in the Python ML pipeline — can be traced directly to one or more of the pain points articulated here.

Chapter 3 begins the technical documentation of the project's solution architecture, with the design and implementation of the **LogisticT SQL Server relational database** — the data governance foundation upon which the entire dual-system intelligence platform is built.

CHAPTER 3: Database Architecture & Business Intelligence (SQL)

The data layer of any enterprise-grade logistics intelligence system constitutes its most critical foundation. Without a well-governed, rigorously normalized, and constraint-enforced relational database, the analytical dashboards and predictive models built on top of it would yield unreliable, inconsistent, and ultimately worthless insights. This chapter documents the complete database architecture of the **LogisticT** database — a purpose-built SQL Server implementation designed to serve as the single source of truth for all operational, financial, and geographical data pertaining to the transportation and logistics tracking system.

The architecture is organized around three principal concerns: (1) the structural design of the relational schema, including normalization strategy, referential integrity, and constraint enforcement; (2) the semantic design of the core fact tables that capture the end-to-end lifecycle of every shipment; and (3) the analytical layer, represented by a suite of ten SQL **VIEWS** that translate raw transactional data into actionable business intelligence for executive decision-making, operational monitoring, and financial reporting.

3.1 Relational Database Design & Entity-Relationship Architecture

3.1.1 Schema Overview

The **LogisticT** database was engineered on **Microsoft SQL Server** and comprises **eleven (11) interrelated tables**, each representing a distinct domain entity within the logistics operation. The schema follows **Third Normal Form (3NF)** principles throughout, ensuring that all non-key attributes depend exclusively on the primary key of their respective table, thereby eliminating transitive dependencies and update anomalies.

#	Table Name	Domain Entity & Purpose	Primary Key
1	Customers	End-recipients of shipped parcels; stores customer contact and location data	CustomerID (VARCHAR)
2	O_Hubs	Origin dispatch hubs; records location, state, and capacity of outbound logistics centers	HubID (VARCHAR)
3	D_Hubs	Destination receiving hubs; mirrors O_Hubs schema for inbound logistics centers	HubID (VARCHAR)
4	Vehicles	Fleet registry; associates each vehicle registration with its classification type	VehicleNumber (VARCHAR)
5	Routes	Named transportation corridors connecting origin and destination hubs	RouteID (INT, IDENTITY)
6	Sellers	Shipper/sender entities; typically commercial clients dispatching goods	SellerID (VARCHAR)
7	DeliveryPersons	Last-mile courier registry; uses auto-incremented numeric ID to enforce uniqueness	DeliveryPersonID (INT, IDENTITY)
8	Parcels	Core financial & booking fact table; the primary anchor for all shipment transactions	TrackingNumber (VARCHAR)
9	Parcels_Details	Operational detail extension of Parcels; stores routing, assignment, and failure data	ParcelID (INT, IDENTITY)
10	Shipment_Logs	Time-series GPS movement log; records every physical location checkpoint per shipment	LogID (INT, IDENTITY)
11	Revenue	Financial reconciliation table; links vehicle maintenance costs to shipment transactions	BillID (INT, IDENTITY)

Table 3.1 — Complete Schema Entity Registry (LogisticT Database)

3.1.2 Entity-Relationship Diagram

The relational structure of the **LogisticT** database is visually represented in the Entity-Relationship (ER) diagram below. The diagram illustrates all **Primary Key** and **Foreign Key** relationships across the eleven tables, depicting how the central fact tables (**Parcels**, **Parcels_Details**) act as the gravitational center of the schema, referencing all surrounding dimension and reference tables.




Figure 3.1 — Entity-Relationship Diagram: LogisticT Database (SQL Server)

3.1.3 Data Integrity, Normalization & Constraint Engineering

A fundamental design principle of this schema is the doctrine that **data quality must be enforced at the database layer**, not deferred to the application or analytics layers. The following mechanisms were implemented to guarantee referential integrity, domain correctness, and consistency throughout the data pipeline.

(A) Surrogate Key Strategy & Normalization

To comply rigorously with **First Normal Form (1NF)** and **Second Normal Form (2NF)**, all entities that could theoretically be identified by human-readable string values were replaced with unique numerical identifiers as their **Primary Keys**. The most critical example is the **DeliveryPersons** table, where string-based names — which are inherently ambiguous, prone to typographical variation, and non-unique — were replaced with a system-generated **INT IDENTITY(1,1) Primary Key** named **DeliveryPersonID**. This numerical key serves as the sole reference point across all related tables, ensuring that personnel assignments in **Parcels_Details** are always unambiguous and that any update to an agent's name propagates correctly via a single record update without causing anomalies across the dependent tables.

(B) Domain Constraints via CHECK and DEFAULT

Multiple **CHECK constraints** and **DEFAULT values** were defined at the DDL level to enforce business rules before any record enters the database. This approach ensures that even if the front-end application or data ingestion pipeline submits invalid data, the database engine itself will reject the transaction.

Constraint Type	SQL Implementation	Business Rule Enforced
CHECK — Non-Negative Fee	ShippingFee >= 0	Prevents the recording of a negative shipping charge, which would corrupt revenue calculations in Power BI.
CHECK — Status Enumeration	Status IN ('Delayed','OnTime','Early','Failed','Pending')	Restricts the shipment status field to a controlled vocabulary, preventing free-text values that would break dashboard filters.
CHECK — Numeric Contact	ISNUMERIC(ContactNumber_C) = 1	Validates that contact numbers contain only numeric characters, enforcing format consistency for operational communications.
DEFAULT — Booking Timestamp	BookingDate DEFAULT GETDATE()	Automatically records the precise booking timestamp if the application layer does not provide one, ensuring no untracked entries.
DEFAULT — Pending Status	Status DEFAULT 'Pending'	Ensures every newly created shipment begins in the 'Pending' state, establishing a consistent lifecycle starting point.
DEFAULT — Zero Maintenance	Maintenance DEFAULT 0	Defaults vehicle maintenance cost to zero in the Revenue table, preventing NULL values that would disrupt profit calculations.
NOT NULL — Critical Fields	CustomerID NOT NULL, VehicleType NOT NULL	Mandates that every parcel is linked to a customer and every vehicle has a classified type, preventing orphaned records.
FOREIGN KEY — Referential Integrity	Parcels_Details REFERENCES Parcels(TrackingNumber)	Guarantees that no parcel detail record can exist without a corresponding parent booking record, maintaining chain of custody.

Table 3.2 — Database-Level Constraint Registry

(C) Referential Integrity via Foreign Keys

The schema enforces a comprehensive network of **Foreign Key (FK) constraints** across the eleven tables. The central hub of this referential web is the **TrackingNumber** field, which originates as the **Primary Key** of the **Parcels** table and propagates as a **Foreign Key** into **Parcels_Details**, **Shipment_Logs**, and **Revenue**. This design pattern ensures that all operational, geographical, and financial records for any shipment are irrevocably linked to a single parent booking record, making the system inherently auditable and traceable.

3.2 The Core Fact Tables: Shipment Tracking Architecture

While all eleven tables contribute to the completeness of the data model, three tables form the operational backbone of the entire logistics tracking system: **Parcels**, **Parcels_Details**, and **Shipment_Logs**. Understanding the deliberate separation of concerns between these tables is essential to appreciating the analytical power they unlock.

3.2.1 The Parcels Table — Financial & Booking Metadata

The **Parcels** table is the primary **fact table** of the schema. Its purpose is to serve as the master booking record for every shipment entered into the system. Critically, it handles all **financial and temporal metadata** at the transaction level, including:

- **TrackingNumber** — The globally unique identifier for every shipment (Primary Key).
- **ShippingFee** — The gross revenue generated by the shipment, protected by a non-negative CHECK constraint.
- **BookingDate** — The precise timestamp of the shipment booking, automatically populated via DEFAULT GETDATE().
- **PromisedDate** — The contractual SLA delivery commitment date, forming the baseline for breach detection.
- **DeliveryDate** — The actual delivery timestamp, compared against PromisedDate to calculate performance.
- **Status** — The high-level lifecycle status of the shipment, restricted to five enumerated values by CHECK constraint.
- **WeatherCondition** — An external factor record, enabling correlation analysis between environmental conditions and delivery outcomes.

3.2.2 The Parcels_Details Table — Operational & Route Assignment

The **Parcels_Details** table extends each booking record with all **operational, geographical, and assignment data** that governs how the shipment is physically executed. It answers the operational questions: *Who delivered it? With what vehicle? On which route? From which hub to which hub?* Key attributes include:

- **SourceHubID / DestinationHubID** — Foreign Keys linking to O_Hubs and D_Hubs respectively, establishing the physical origin-to-destination corridor.
- **DeliveryPersonID** — Foreign Key referencing DeliveryPersons, enabling courier-level performance attribution.
- **VehicleID** — Foreign Key referencing the Vehicles table, enabling fleet utilization and cost analysis.
- **RouteID** — Foreign Key referencing the Routes table, enabling route-level efficiency benchmarking.
- **Material_Type** — Classifies the cargo category, enabling segment-level analytics (e.g., hazardous goods vs. standard parcels).
- **Failure_Reason** — A nullable text field that captures the root cause of a failed delivery attempt, feeding directly into the **vw_DeliveryAttempts** diagnostic view.

3.2.3 The Shipment_Logs Table — Time-Series GPS Tracking

The **Shipment_Logs** table constitutes the **real-time event stream** of the system. Unlike the static record-per-shipment model of the Parcels tables, Shipment_Logs employs a **one-to-many relationship** with Parcels, capturing every physical location checkpoint that a shipment passes through during transit. This design generates a **time-series audit trail** that enables both retrospective route analysis and real-time operational alerts. Key attributes include:

- **LogID** — Auto-incremented surrogate key ensuring log immutability.
- **Log_Timestamp** — Precise datetime of the GPS checkpoint, defaulting to GETDATE() at insertion.
- **Location** — Human-readable or GPS-coordinate-based checkpoint description.
- **GpsProvider** — Tracks which GPS telematics provider recorded the checkpoint, enabling data source auditing.

Architectural Design Note: The Separation of Concerns Pattern

The deliberate separation of the **Parcels** (financial/contractual layer), **Parcels_Details** (operational/assignment layer), and **Shipment_Logs** (real-time movement layer) mirrors the architectural principle of separation of concerns. Each table can be queried and updated independently without locking or degrading the others, ensuring optimal performance under high-concurrency logistics operations. This three-tier fact structure is the central nervous system of the entire intelligence platform.

3.3 Business Intelligence Queries: The Ten Executive Views

A core deliverable of the database architecture phase was the development of **ten SQL VIEWS** that pre-compute the most critical analytical queries required by operational managers, financial controllers, and logistics engineers. Rather than exposing raw transactional tables to downstream systems — which would require analysts to possess deep SQL expertise and would result in inconsistent query logic — the **VIEW layer** acts as a semantic abstraction that encapsulates complex multi-table **JOIN** logic, aggregations, and window functions into reusable, named query objects.

These views are consumed directly by **Power BI** as data sources for executive dashboards, and their output structures informed the feature engineering pipeline used by the **Python machine learning models** described in subsequent chapters.

#	View Name	Business Purpose	Core Tables Used
1	vw_ShipmentMaster	End-to-end shipment overview from Sender to Receiver	Parcels, Parcels_Details, Sellers, Customers, O_Hubs, D_Hubs
2	vw_LatestShipmentStatus	Real-time last known location per shipment	Shipment_Logs (ROW_NUMBER window fn.)
3	vw_TrackingTimeline	Full chronological movement history per parcel	Parcels, Shipment_Logs
4	vw_HubDailyWorkload	Daily outbound parcel count per origin hub	Parcels, Parcels_Details
5	vw_CourierPerformance	Courier success rate: OnTime+Early vs. total assigned	DeliveryPersons, Parcels_Details
6	vw_DeliveryAttempts	Root-cause failure frequency aggregation by reason	Parcels_Details (failed records)
7	vw_LateDeliveryReport	SLA breach report with exact delay hours calculation	Parcels, Parcels_Details
8	vw_StuckShipments	Operational alert: shipments with no movement for 48+ hours	Parcels, Shipment_Logs
9	vw_PODValidation	Proof-of-delivery audit: GPS provider presence check	Parcels, Shipment_Logs
10	vw_HubRevenueMonthly	Monthly hub-level P&L: Net Profit = Fees minus Maintenance	Parcels, Parcels_Details, Revenue

Table 3.3 — Business Intelligence View Registry (LogisticT Database)

The following subsections provide a detailed technical and business analysis of the five most strategically significant views in the suite.

3.3.1 vw_ShipmentMaster — The End-to-End Operational Ledger

Business Justification

The single most frequent query in any logistics operation is: *"What is the current status of shipment X, who sent it, who is receiving it, and between which hubs is it travelling?"* Without a consolidated view, answering this question requires a minimum of five separate table lookups and complex manual JOIN logic. The **vw_ShipmentMaster** view eliminates this friction entirely, providing a single, denormalized projection that integrates six underlying tables into one coherent record per parcel.

Technical Implementation

```
CREATE VIEW vw_ShipmentMaster AS
SELECT
    pd.ParcelID,
    pd.TrackingNumber,
    s.Name          AS SenderName,
    c.Name          AS ReceiverName,
    oh.Location     AS Origin,
    dh.Location     AS Destination,
    p.Status,
    p.DeliveryDate
FROM Parcels_Details pd
JOIN Parcels      p  ON pd.TrackingNumber = p.TrackingNumber
JOIN Sellers      s  ON pd.SellerID       = s.SellerID
JOIN Customers    c  ON p.CustomerID      = c.CustomerID
JOIN O_Hubs       oh ON pd.SourceHubID    = oh.HubID
JOIN D_Hubs       dh ON pd.DestinationHubID = dh.HubID;
```

Analytical Value

This view forms the primary data source for the Master Shipment Tracker page within the Power BI dashboard. By pre-joining six tables, it enables Power BI to render an interactive shipment register without placing any JOIN computation burden on the reporting layer. It also provides the ground truth dataset used in the Python model to engineer features such as route distance and hub pairing frequency.

3.3.2 vw_LateDeliveryReport — SLA Breach Detection Engine

Business Justification

Service Level Agreement (SLA) compliance is the primary contractual obligation in commercial logistics. Every hour a parcel exceeds its **PromisedDate** represents a breach of the service contract, a potential financial penalty, and a reputational risk. The **vw_LateDeliveryReport** view operationalizes SLA monitoring by calculating the precise **DelayHours** for every shipment that has been delivered beyond its commitment window.

Technical Implementation

```
CREATE VIEW vw_LateDeliveryReport AS
SELECT
    p.TrackingNumber,
    pd.SourceHubID      AS OriginHub,
    pd.DestinationHubID AS DestinationHub,
    p.PromisedDate,
    p.DeliveryDate      AS ActualDate,
    DATEDIFF(HOUR, p.PromisedDate, p.DeliveryDate) AS DelayHours
FROM Parcels p
JOIN Parcels_Details pd ON p.TrackingNumber = pd.TrackingNumber
WHERE p.Status = 'Delayed'
AND p.DeliveryDate > p.PromisedDate;
```

Analytical Value

The **DATEDIFF(HOUR, PromisedDate, DeliveryDate)** expression produces a precise integer delay value that is consumed by Power BI to render an SLA breach severity chart. Additionally, this **DelayHours** metric serves as the primary **target variable (label)** for the supervised regression and classification models developed in the Python machine learning chapter, making this view the critical bridge between the SQL operational layer and the predictive analytics layer.

3.3.3 vw_DeliveryAttempts — Root Cause Failure Analysis

Business Justification

When a delivery attempt fails, the organization incurs re-delivery costs, customer dissatisfaction, and potential SLA penalties. To reduce the failure rate, management must understand the **systemic causes** of failure rather than treating each incident in isolation. The **vw_DeliveryAttempts** view serves as an automated root cause aggregation engine, counting the frequency of each distinct **Failure_Reason** across all failed delivery records.

Technical Implementation

```
CREATE VIEW vw_DeliveryAttempts AS
SELECT
    Failure_Reason,
    COUNT(TrackingNumber) AS FailedAttempts
FROM Parcels_Details
WHERE Status = 'Failed'
AND Failure_Reason IS NOT NULL
GROUP BY Failure_Reason;
```

Analytical Value

This view powers the **Failure Root Cause Analysis** visual in Power BI, enabling management to identify the leading causes of delivery failures (e.g., "*Traffic Congestion*", "*Vehicle Breakdown*", "*Customer Refused Delivery*") through a Pareto chart. The inclusion of sample data records for **TRK_FAIL_001** (Failure Reason: Customer Refused Delivery) and **TRK_FAIL_002** (Failure Reason: Wrong Address) immediately validates the view's logic during testing and demonstrates two distinct failure archetypes for analytical modeling.

3.3.4 vw_HubRevenueMonthly — Financial Health & Profitability Dashboard

Business Justification

Understanding which operational hubs are profitable and which are generating losses is essential for capital allocation and network optimization decisions. The **vw_HubRevenueMonthly** view computes a monthly Profit & Loss summary at the hub level by calculating **Net Profit** as the difference between total shipping fee revenue and total vehicle maintenance expenditure.

Technical Implementation

```
CREATE VIEW vw_HubRevenueMonthly AS
SELECT
    FORMAT(p.BookingDate, 'yyyy-MM')      AS YearMonth,
    pd.SourceHubID                        AS Hub,
    COUNT(p.TrackingNumber)              AS TotalShipments,
    SUM(p.ShippingFee)                   AS TotalShippingFees,
    SUM(r.Maintenance)                   AS TotalCosts,
    SUM(p.ShippingFee - ISNULL(r.Maintenance, 0)) AS NetProfit
FROM Parcels p
JOIN Parcels_Details pd ON p.TrackingNumber = pd.TrackingNumber
JOIN Revenue r          ON p.TrackingNumber = r.TrackingNumber
GROUP BY FORMAT(p.BookingDate, 'yyyy-MM'), pd.SourceHubID;
```

Analytical Value

Several design decisions in this view deserve architectural recognition. The **FORMAT(BookingDate, 'yyyy-MM')** expression creates a consistent monthly time-series grain that Power BI can interpret as a date hierarchy for trend analysis. The use of **ISNULL(r.Maintenance, 0)** prevents NULL propagation in the arithmetic, which would otherwise render the entire **NetProfit** calculation null for any shipment with a missing maintenance record — an issue that would otherwise silently corrupt all financial reporting. This view directly feeds the **Hub Financial Performance** KPI cards and trend charts in the Power BI executive dashboard.

3.3.5 vw_StuckShipments — Real-Time Operational Alert System

Business Justification

In active logistics networks, one of the most operationally dangerous failure modes is a shipment that has entered the **'Pending'** or in-transit state but has ceased generating movement logs — an indicator of a potential vehicle breakdown, route blockage, or system data ingestion failure. The **vw_StuckShipments** view functions as a **real-time operational alert**, automatically identifying all shipments whose last recorded GPS checkpoint timestamp is more than **48 hours** in the past relative to the current server time.

Technical Implementation

```
CREATE VIEW vw_StuckShipments AS
SELECT
    p.TrackingNumber,
    ls.Location,
    ls.Log_Timestamp
FROM Parcels p
JOIN Shipment_Logs ls ON p.TrackingNumber = ls.TrackingNumber
WHERE p.Status = 'Pending'
AND DATEDIFF(HOUR, ls.Log_Timestamp, GETDATE()) > 48;
```


Analytical Value

The critical technical feature of this view is the use of **GETDATE()** within the view's **WHERE clause**. Unlike static threshold comparisons, this approach means the view's result set *evolves dynamically with the passage of real time*. Every time the view is queried — whether by Power BI on a scheduled refresh or by a monitoring application — it re-evaluates the 48-hour threshold against the current server timestamp, effectively making it a **live operational dashboard**. The sample record **TRK_STUCK_001** — last logged at *Stopped at Vellore Toll* — was specifically inserted into the test dataset to validate this alert logic.

3.4 Supporting Views: Operational Completeness

Beyond the five primary views analysed in detail above, the remaining five views in the suite contribute to operational completeness and data governance.

- **vw_LatestShipmentStatus** — Employs a **ROW_NUMBER() OVER(PARTITION BY TrackingNumber ORDER BY Log_Timestamp DESC)** window function to isolate the most recent location checkpoint for each shipment, providing a real-time position feed without requiring the consumer to resolve timestamp ordering manually.
- **vw_TrackingTimeline** — Produces a chronologically ordered event log for each parcel by joining **Parcels** with all associated **Shipment_Logs** records, enabling the end-customer-facing parcel tracking interface to render a step-by-step delivery history.
- **vw_HubDailyWorkload** — Aggregates daily outbound parcel counts per origin hub by casting the **BookingDate** to a **DATE** grain and grouping on **SourceHubID**, enabling operations managers to monitor hub throughput against capacity limits in real time.
- **vw_CourierPerformance** — Computes each delivery agent's **success rate** by summing all records where **Status** is 'OnTime' or 'Early' against total assigned deliveries, providing the HR and Operations departments with an objective performance scorecard.
- **vw_PODValidation** — Audits the completeness of **Proof of Delivery (POD)** records by verifying that a **GpsProvider** value is present in the most recent **Shipment_Log** entry for all completed shipments, ensuring compliance with contractual documentation requirements.

3.5 Chapter Conclusion: The Data Foundation for Analytics & Intelligence

This chapter has documented the complete database architecture of the **LogisticT** SQL Server system — from the structural foundations of the relational schema through to the analytical abstractions provided by the Business Intelligence VIEW layer. The key engineering outcomes of this phase are summarized as follows:

- **Structural Integrity:** The eleven-table schema, designed to Third Normal Form, eliminates data redundancy and ensures that all updates propagate correctly through the relational hierarchy without anomalies.
- **Constraint-Enforced Quality:** Eight distinct constraint types — including **CHECK**, **DEFAULT**, **NOT NULL**, and **FOREIGN KEY** constraints — enforce business rules at the database engine level, guaranteeing that the data entering the analytical pipeline is structurally and semantically correct.
- **Three-Tier Fact Architecture:** The separation of **Parcels** (financial), **Parcels_Details** (operational), and **Shipment_Logs** (temporal/spatial) into distinct tables mirrors enterprise data warehouse design patterns, enabling independent scaling and querying of each concern.

- **Semantic Analytics Layer:** The ten SQL VIEWS encapsulate complex multi-table JOIN and aggregation logic into named, reusable query objects, dramatically reducing the complexity burden on both the Power BI reporting layer and the Python data ingestion pipeline.

Pathway to the Next Chapters

Chapter 4 — Inbound Analytics (Power BI): The ten views documented in this chapter serve as the direct data connectors for Power BI. Views such as `vw_ShipmentMaster`, `vw_HubRevenueMonthly`, and `vw_CourierPerformance` will be consumed by the Power BI data model to populate KPI cards, trend charts, and operational drill-through reports across the executive dashboard.

Chapter 5 — Predictive Modeling (Python / ML): The clean, structured output of `vw_LateDeliveryReport` — specifically the `DelayHours` metric — will be ingested by the Python pipeline as the primary **target variable** for training delay prediction models. Categorical features such as `WeatherCondition`, `Material_Type`, `RouteID`, and `SourceHubID` — all accessible via the structured schema — will undergo feature encoding and be used as input predictors. The constraint-enforced data quality documented in this chapter ensures that the ML training dataset is free of NULL contamination and domain violations before any model is fitted.

In summary, the **LogisticT** SQL Server database is not merely a data store — it is an actively governed, constraint-enforced, analytically optimized data platform. By investing in the rigour of its relational design and constraint architecture at the foundation layer, the system guarantees that every dashboard visualization and every machine learning prediction built upon it rests on a bedrock of clean, consistent, and trustworthy data.

CHAPTER 4: Inbound Analytics & Warehouse Optimization

A Power BI Intelligence Study on Inventory Risk, Zone Efficiency & Supply Chain Continuity

This chapter presents a comprehensive analytical investigation into the inbound dimension of the dual-system supply chain architecture deployed in this project. While Chapter 3 addressed outbound and last-mile delivery intelligence — powered by Python, SQL, and machine learning — the present chapter directs the analytical lens inward, examining the warehouse environment that forms the **operational foundation** upon which all outbound logistics ultimately depend.

The analysis herein is conducted entirely through a **Power BI enterprise intelligence dashboard** built over an internal warehouse dataset of **3,204 SKUs** spanning five product categories and four warehouse zones. The findings are structured across three analytical pillars: **financial impact assessment**, **inventory risk profiling**, and **operational zone benchmarking**. The chapter concludes with a critical synthesis — The Bridge — which formally connects inbound warehouse performance to outbound delivery reliability, establishing the logical necessity of warehouse optimization as a prerequisite for predictive dispatch intelligence.

4.1 Executive Overview & Financial Impact Assessment

4.1.1 Dataset Scope & Operational Scale

The warehouse dataset underpinning this analysis encompasses a total of **3,204 Stock-Keeping Units (SKUs)**, distributed across **five product categories** — Automotive, Pharma, Electronics, Groceries, and Apparel — and managed across **four discrete warehouse zones** (Zone A, B, C, and D). This breadth of inventory scope provides a statistically robust foundation for cross-category risk profiling and zone-level benchmarking.

The Power BI dashboard aggregates and visualizes this multi-dimensional dataset to expose systemic inefficiencies, quantify financial exposure, and surface actionable intelligence for warehouse managers and supply chain executives.

4.1.2 The Fulfillment Gap: Current Performance vs. Industry Benchmark

The most strategically significant finding of this executive overview is the **measurable gap between current warehouse fulfillment performance and established industry standards**. The current Order Fulfillment Rate stands at **84.97%** — a figure that, while operationally active, represents a **10.03 percentage-point deficit** relative to the **95% industry benchmark** recognized across global logistics and supply chain management bodies.

This gap is not merely a performance statistic — it is a direct proxy for **unmet customer demand, unrecovered revenue, and systemic operational fragility**. The financial translation of this gap is stark and demands executive attention.

Metric	Current State	Target / Benchmark
Order Fulfillment Rate	84.97%	95.00% (Industry Standard)
Fulfillment Gap	-10.03 percentage points	0.00% (Zero Gap)
Total Inventory Value (Current)	\$89.57 Million	—
Potential Monthly Revenue (100% Rate)	\$257.80 Million	Unlocked at Full Fulfillment
Revenue Currently at Risk	~\$38.67 Million / Month*	Estimated from 15.03% gap

** Estimated proportional revenue at risk based on the 15.03% fulfilment shortfall applied against total monthly revenue potential.*

Table 4.1 encapsulates the magnitude of the financial opportunity available to the organization. The current inventory base of **\$89.57 million** represents the capital already deployed within the warehouse network. However, the gap in fulfillment efficiency means that a potential monthly revenue of **\$257.80 million** — representing a ceiling unlocked only at a 100% fulfillment rate — remains partially inaccessible due to structural inventory and operational inefficiencies. Closing even half of the current gap would yield measurable top-line revenue growth without capital reinvestment.

4.2 Inventory Risk & Stockout Analysis

4.2.1 System-Wide Critical Risk Exposure

Across the full portfolio of 3,204 SKUs, the Power BI dashboard identifies **237 SKUs — representing 7.4% of total inventory** — currently positioned **below their designated reorder point**. A reorder point is the inventory threshold at which procurement must be triggered to prevent stockout conditions, calculated as a function of lead time and daily demand variability. Any SKU below this threshold is operating in a **critical risk corridor**: one where a single demand spike or supply disruption can instantly translate into a stockout, a missed customer order, and an unrecoverable service failure.

The operational implication is unambiguous: **immediate procurement intervention** is required for these 237 SKUs. Deferral of replenishment action on this cohort represents a compounding risk — each day of inaction widens the exposure window and increases the probability of multi-SKU concurrent stockout events, which have disproportionate impacts on fulfillment rates and customer satisfaction scores.

CRITICAL ALERT — 237 SKUs (7.4% of Portfolio) Below Reorder Point

Immediate procurement escalation is required. These items present a direct and quantifiable risk to the system-wide Order Fulfillment Rate and must be prioritized in the next procurement cycle.

4.2.2 Category-Level Risk Profiling

While system-level indicators provide a macro-view of inventory health, a category-level disaggregation reveals the **asymmetric distribution of risk and reward** across product families. Two categories warrant elevated executive attention: Automotive and Pharma.

Category	Stockout Count	Critical SKUs	Inventory Value	Potential Monthly Revenue
Automotive	3,110	51	—	\$54.4M / month
Pharma	High	Elevated	\$19.14M	High-margin + compliance
Electronics	Moderate	Moderate	—	Moderate
Groceries	Moderate	Low	—	High volume, lower margin
Apparel	Low-Moderate	Low	—	Seasonal sensitivity

Table 4.2 — Category-Level Risk and Revenue Profile

4.2.3 Automotive Category — High Risk, High Reward

The **Automotive category** presents the most pronounced risk-reward duality in the portfolio. It records the highest stockout volume in the dataset at **3,110 stockout occurrences** in the measurement period — a figure that represents not merely a procurement failure but a pattern of systemic demand-supply misalignment. Concurrently,

51 SKUs within this category are classified as **critical products**, meaning their absence triggers cascading production or service disruption for downstream customers.

Yet this risk exposure is paired with the highest revenue opportunity in the portfolio: **\$54.4 million per month** in potential recoverable revenue. This positions the Automotive category as the **highest-leverage intervention point** within the warehouse inventory strategy. A targeted replenishment and safety-stock recalibration program for Automotive SKUs — particularly those 51 critical products — would yield the greatest marginal return per dollar of procurement investment.

4.2.4 Pharma Category — Dual Financial & Compliance Risk

The **Pharmaceutical category** introduces a dimension of risk that transcends financial exposure. With an inventory holding value of **\$19.14 million** — the highest of any category — and a handling cost of **\$2.849 per unit** (the highest across all five categories), the Pharma segment generates a disproportionate carrying cost burden relative to its throughput volume.

More critically, any stockout event within the Pharmaceutical category generates a **dual-risk impact**: a financial penalty in the form of lost sales and emergency procurement premiums, compounded by **healthcare compliance and regulatory risk**. Pharmaceutical inventory management is governed by regulatory frameworks that mandate product availability, traceability, and chain-of-custody continuity. A stockout in this category is not merely a missed shipment — it is a potential compliance breach, liability exposure, and patient safety concern. The Power BI analysis therefore identifies Pharma as a **priority category for safety stock augmentation** and procurement lead-time reduction.

4.3 Operational Efficiency & Zone Performance Analysis

4.3.1 Framework: The Zone KPI Model

Warehouse operational efficiency in this system is evaluated through a composite **Zone KPI Score** — a weighted performance index that integrates picking time, layout efficiency, and order fulfillment rate at the zone level. This metric provides a holistic, quantitative basis for comparing operational performance across the four warehouse zones (A, B, C, and D), and serves as the primary benchmarking instrument for warehouse layout optimization decisions.

The KPI score is normalized on a 0.000–1.000 scale, where higher values indicate superior operational performance. **Picking time** — the duration in seconds required to locate, retrieve, and stage a SKU for dispatch — is the primary operational driver embedded within this score. A lower picking time directly correlates with faster order processing, higher throughput, and reduced truck loading delays, all of which propagate downstream into last-mile delivery performance.

4.3.2 Zone Performance Comparison

Zone	KPI Score	Avg. Picking Time	Performance Tier	Priority
Zone A	0.606 ★ Benchmark	94 seconds (Fastest)	Tier 1 — Elite	Model Zone
Zone B	TBD (Sub-benchmark)	> 94 seconds	Tier 2 — Adequate	Optimization Target
Zone C	TBD (Sub-benchmark)	> 94 seconds	Tier 2 — Adequate	Optimization Target
Zone D	TBD (Sub-benchmark)	> 94 seconds	Tier 2/3	Urgent Restructure

Table 4.3 — Warehouse Zone Performance Matrix

4.3.3 Zone A — The Operational Benchmark

Zone A emerges from the Power BI zone analysis as the definitive operational benchmark within the warehouse network. It achieves a **KPI score of 0.606** — the highest recorded across all four zones — and a mean picking time of just **94 seconds per SKU**. This combination of high KPI attainment and minimal picking latency establishes Zone A as the proof-of-concept for effective warehouse layout design.

The performance characteristics of Zone A are attributable to a combination of factors discernible within the dataset: high layout efficiency scores, logical SKU proximity and slotting strategy (high-velocity items positioned near dispatch staging areas), and effective personnel routing. The operational intelligence embedded within Zone A's layout strategy is **a replicable organizational asset**, not an anomaly.

Strategic Recommendation:

The layout optimization principles embedded in Zone A — including SKU slotting logic, picking route design, and staging efficiency — should be formally documented and systematically replicated across Zones B, C, and D. A phased zone restructuring program, benchmarked against Zone A KPI metrics, is recommended as a top-priority operational initiative in the warehouse improvement roadmap.

4.3.4 Zones B, C, and D — Optimization Imperatives

In contrast to Zone A, the remaining three zones operate at sub-benchmark KPI levels, with correspondingly higher average picking times. The performance gap between these zones and Zone A is not a marginal variance — it is a **structural inefficiency** that accumulates across thousands of daily pick operations to produce measurable delays in order processing, truck loading, and dispatch scheduling.

The operational priority hierarchy is clear: any zone recording a picking time materially above 94 seconds is a zone generating unnecessary latency in the order fulfillment pipeline. This latency does not remain contained within the warehouse — it propagates outward into the transportation network, manifesting as delayed dispatch windows, tighter delivery time buffers, and elevated last-mile delay probabilities. The analysis in Section 4.4 formalizes this causal linkage.

4.4 The Bridge: How Warehouse Operations Drive Last-Mile Delivery Outcomes

"Delivery delays do not begin on the road. They begin in the warehouse."

— Chapter 4 Synthesis Finding

4.4.1 The Causal Chain: From Warehouse Floor to Last-Mile Performance

A foundational assumption embedded in many logistics intelligence systems is that delivery performance is primarily a function of **outbound variables**: road conditions, carrier capacity, route optimization, and traffic patterns. This project challenges that assumption with empirical evidence drawn from the internal warehouse dataset. The Power BI analytics presented in Sections 4.1 through 4.3 systematically demonstrate that the supply chain conditions governing last-mile delivery outcomes are established — and frequently compromised — long before a vehicle departs the distribution center.

The causal chain is direct, quantifiable, and operationally actionable. It begins in the warehouse and terminates at the customer's doorstep. Understanding this chain is not merely an academic exercise — it is the analytical justification for the **dual-system architecture** that defines this project's methodology.

4.4.2 Causal Link 1: Picking Time Inefficiency → Delayed Truck Dispatch

Every second of excess picking time accumulated within Zones B, C, and D relative to Zone A's 94-second benchmark represents a quantifiable delay introduced into the order preparation pipeline. At scale, across hundreds or thousands of daily pick operations, these seconds aggregate into minutes — and minutes aggregate into **missed dispatch windows**.

Truck dispatch schedules operate within rigid time windows constrained by driver hours of service regulations, route optimization requirements, and last-mile delivery commitments. A warehouse that consistently fails to complete order picking and staging operations within the required preparation window **forces one of two equally damaging outcomes**: either the dispatch truck departs with an incomplete load — reducing load utilization and revenue per trip — or the truck departs late, compressing the available delivery window and escalating the probability of a last-mile delay.

- **Operational Implication:** Reducing picking time in Zones B, C, D toward the Zone A benchmark of 94 seconds is not a warehouse-level efficiency goal in isolation — it is a **dispatch reliability intervention** that directly protects last-mile delivery performance.

4.4.3 Causal Link 2: Stockouts & Fulfillment Failures → Order Fragmentation & Re-dispatch

The 237 SKUs currently below reorder point — and the 3,110 Automotive stockout events recorded in the measurement period — each represent an instance where a customer order could not be fulfilled from existing inventory. The downstream consequences of these events are not limited to the immediate lost sale; they generate a cascade of operational disruptions that ripple through the outbound delivery network.

When a stockout occurs, order management systems typically execute one or more of the following responses: partial shipment (fragmenting the order across multiple deliveries), order holding (awaiting replenishment, introducing processing lag), emergency procurement (at premium cost), or order cancellation (the costliest outcome). Each of these responses **adds complexity, cost, and delay** to the outbound delivery operation.

- **Order Fragmentation:** Partial orders require separate dispatch events, doubling or tripling the route complexity for a single customer commitment.
- **Processing Lag:** Held orders consume warehouse staging capacity and occupy scheduling slots in the dispatch queue, delaying other orders in the pipeline.
- **Emergency Procurement:** Unplanned replenishment introduces unpredictable lead times that the predictive ETA models in the outbound system cannot account for, degrading forecast accuracy.

4.4.4 Causal Link 3: Fulfillment Rate → Predictive Model Input Validity

The outbound analytics system — detailed in subsequent chapters — deploys **machine learning models for predictive ETA estimation** and dynamic dispatch scheduling. These models are trained on historical order and delivery data and produce forecasts that are only as reliable as the **consistency and integrity of the inputs** they receive from the warehouse system.

An inbound environment characterized by high stockout rates, variable picking times, and sub-optimal zone performance generates **irregular and unpredictable dispatch patterns**. Trucks depart at inconsistent times. Order sizes vary unpredictably. Carrier assignments shift due to last-minute order changes. This operational volatility **degrades the quality of the historical data** on which the predictive ML models depend, reducing their forecast accuracy and undermining the reliability of ETA predictions communicated to customers and carriers.

Conversely, a warehouse operating at or near Zone A's KPI benchmark — with consistent picking times, high fulfillment rates, and zero stockout events — produces **stable, predictable dispatch patterns** that the outbound ML models can leverage for high-confidence ETA estimation. **Warehouse optimization is therefore not merely an operational improvement; it is a prerequisite for reliable predictive intelligence in the outbound system.**

4.4.5 The Holistic System: Inbound and Outbound as a Unified Intelligence Architecture

This chapter has demonstrated that the inbound analytics layer of the dual-system architecture is not a standalone analytical module — it is the **root system** upon which the entire supply chain intelligence tree depends. The financial exposure quantified in Section 4.1, the inventory risks profiled in Section 4.2, and the zone inefficiencies benchmarked in Section 4.3 are not warehouse-level issues with warehouse-level consequences. They are **system-wide conditions** with measurable, traceable impacts on the outbound delivery performance analyzed in the chapters that follow.

The logical architecture of this project — inbound Power BI analytics informing outbound Python and SQL intelligence — reflects a **supply chain truth** that operational data consistently validates: **you cannot optimize what you deliver until you optimize how you prepare it**. The warehouse is not the end of the supply chain process; it is its engine room.

Inbound Condition (This Chapter)	Direct Outbound Impact	System Consequence
High Picking Time (Zones B/C/D)	Late truck dispatch windows	Compressed delivery buffer → higher last-mile delay risk
237 SKUs Below Reorder Point	Order fragmentation & holds	Increased re-dispatch events & customer SLA breaches
84.97% Fulfillment Rate	Volatile dispatch patterns	Reduced ML model accuracy for ETA prediction
3,110 Automotive Stockouts	Emergency procurement lag	Unpredictable carrier scheduling disruptions
Zone A Benchmark Achieved	Consistent dispatch timing	High-confidence ML ETA predictions & SLA compliance

Table 4.4 — Inbound-to-Outbound Impact Matrix: The Bridge Between Warehouse and Last-Mile Performance

In conclusion, the analytical work presented in this chapter establishes that **warehouse optimization is the mandatory first step** in any credible effort to improve last-mile delivery performance. The stockout mitigation strategies, zone layout improvements, and fulfillment rate recovery actions recommended herein are not supplementary to the outbound intelligence system — they are its **preconditions**. The predictive dispatch models, route optimization algorithms, and delivery ETA frameworks presented in subsequent chapters assume a warehouse baseline that can only be achieved through the systematic application of the insights generated by the Power BI inbound analytics system described above.

Chapter 5 will transition from the warehouse floor to the road network, engaging the Python- and SQL-driven outbound analytics layer — with the full understanding that its performance ceiling is determined, in large part, by the inbound operational excellence established here.

Chapter 5: Predictive Analytics & Data Mining (Python)

This chapter documents the full analytical pipeline applied to the outbound logistics dataset using Python. The workflow progresses systematically from raw data ingestion and rigorous engineering, through advanced exploratory analysis and unsupervised pattern discovery, to the deployment of a supervised machine learning classifier for proactive delay prediction. Each methodological step is grounded in established data science principles and is directly motivated by the operational challenges identified in the project's problem statement.

5.1 Data Engineering & Feature Construction

Before any analytical technique can be applied, a logistics dataset must be transformed from a collection of raw operational records into a coherent, analytically tractable structure. This section describes the engineering decisions that shaped the quality and richness of the feature space used in all downstream analyses.

5.1.1 Geospatial Distance Computation: The Haversine Formula

A central operational variable in any logistics analysis is the physical distance between the origin and destination of each shipment. The dataset provides GPS coordinates for each origin hub and destination location in the form of latitude–longitude pairs, stored in the **Org_lat_lon** and **Des_lat_lon** columns respectively. However, a non-trivial proportion of records contained missing or erroneous values in the **TRANSPORTATION_DISTANCE_IN_KM** field. Rather than discarding these records — which would introduce selection bias — the project team recovered the missing distance values through a principled geospatial computation.

The **Haversine Formula** was implemented to compute the great-circle distance — the shortest path between two points on the surface of a sphere — from the raw GPS coordinates. This formula is the industry standard for computing geographic distances over the Earth's curved surface, offering a practical balance between computational efficiency and geodetic precision sufficient for road-freight analytics. Mathematically, the formula is expressed as:

Haversine Formula — Great-Circle Distance

$d = 2R \cdot \arctan2(\sqrt{a}, \sqrt{1-a})$ where: $a = \sin^2(\Delta\phi/2) + \cos(\phi_1) \cdot \cos(\phi_2) \cdot \sin^2(\Delta\lambda/2)$ ϕ_1, ϕ_2 = origin and destination latitudes (radians) $\Delta\phi, \Delta\lambda$ = differences in latitude and longitude R = Earth's mean radius $\approx 6,371$ km

In the project's implementation, the coordinate strings were parsed and converted to radians before applying the formula. This imputation step was performed exclusively on records where the distance field was null and valid GPS coordinates existed, ensuring that originally populated distance values were not overwritten. The result was a complete, geospatially consistent distance column that served as the foundation for multiple downstream features, including route segmentation, speed computation, and cluster profiling.

Business Value of the Haversine Implementation: By recovering missing distance data through geospatial computation rather than mean/median imputation, the team preserved the true heterogeneity of the route network. Imputing with a global mean would have systematically misrepresented short-haul and long-haul routes alike, corrupting speed calculations and bias-ing any model that used distance as a feature. The Haversine-based imputation thus directly improved the integrity of every analytical output produced in this project.

5.1.2 Derived Operational Features

Beyond distance recovery, a structured feature engineering process created thirteen new operational variables designed to capture the temporal, categorical, and performance dimensions of each shipment. These features were constructed exclusively from information available at booking time, a design constraint that prevents data

leakage and ensures all downstream predictive models reflect realistic operational conditions. The principal engineered features are summarised below:

Feature Name	Source	Operational Purpose
target_delay	actual_eta > Planned_ETA	Binary classification target (1 = delayed, 0 = on-time)
delay_days	ETA gap in days	Continuous severity measure for delayed shipments
log_distance	log1p(TRANSPORT_DIST_KM)	Normalised distance for modelling (right-skew correction)
avg_speed	Distance ÷ Duration (km/h)	Physics-based quality signal; outliers flag GPS failures
duration_hr	trip_end_date – trip_start_date	Actual elapsed travel time in hours
pickup_hour	Hour of BookingID_Date	Captures time-of-day operational rhythms
day_of_week / is_weekend	DayOfWeek of BookingID_Date	Identifies weekend operational constraints
booking_month	Month of BookingID_Date	Enables seasonal pattern detection
time_period	Binned pickup_hour	Categorical: Morning / Afternoon / Evening / Night
veh_cat	Encoded vehicleType	Reduced-cardinality vehicle class for modelling
mat_cat	Encoded material type	Commodity category for association rule mining
origin_state / dest_state	Extracted from location codes	Geographic provenance for regional pattern analysis
route	Origin prefix + Dest prefix	Unique origin-destination corridor identifier

5.2 Data Cleaning & Quality Assurance

A rigorous multi-stage cleaning protocol was applied before any analytical technique was executed. Cleaning was governed by the principle that each removal or imputation decision must be justified by a documented business rule rather than applied wholesale. The following sequential validation steps were implemented:

- Logical Timeline Validation:** Four business rules were enforced — arrival dates cannot precede booking dates; trip start cannot precede booking; trip end cannot precede start; and all dates must be post-2010. Records violating any rule were removed as they introduce systematic bias into all date-derived features.
- Speed-Based Outlier Removal:** Average speed was computed for every shipment. Records with speeds below 1 km/h or above 120 km/h were flagged as GPS tracking failures or data corruption events and excluded, as they would distort route time estimations and model training.
- Extreme Duration Filter:** Trips exceeding 360 hours (15 days) were removed, as these almost universally represent cases where GPS tracking devices were not switched off after delivery completion — a hardware artefact rather than a legitimate journey.
- Missing Value Audit:** A systematic column-level audit quantified the null rate, unique cardinality, and data type of every field. Records missing values in critical operational identifiers (vehicle type, origin/destination

codes, actual ETA) were removed. Non-critical missing values were handled through targeted imputation strategies including the Haversine formula for distance and median imputation for other numeric fields.

- **Categorical Normalisation:** All free-text categorical columns were standardised to Title Case with leading/trailing whitespace removed, eliminating phantom duplicate categories caused by inconsistent data-entry formatting.

5.3 Advanced Exploratory Data Analysis

Following data preparation, an extensive multi-layered exploratory analysis was conducted to characterise the distribution of key variables, surface relationships between operational factors, and identify the structural patterns underlying delivery delay. The analyses described in this section range from univariate distribution profiling to complex multivariate relationship mapping.

5.3.1 Normalising Skewed Distributions: Log-Transformation

Inspection of the transportation distance distribution revealed a pronounced **positive (right) skew**: the majority of shipments were short-to-medium haul routes (under 500 km), while a long tail of inter-regional routes extended beyond 2,000 km. This distributional asymmetry violates the normality assumptions of linear statistical methods and degrades the performance of distance-sensitive algorithms such as K-Means clustering and Logistic Regression.

The standard statistical remedy for right-skewed, non-negative continuous variables is the **Log-Transformation**. Specifically, the *log1p* transformation — defined as $\log(1 + x)$ — was applied to avoid computational errors at zero and to produce a feature distribution that is approximately Gaussian. The transformed **log_distance** feature was used in all distance-sensitive models, while the original kilometre values were retained for interpretable reporting and business dashboards. Trip duration exhibited the same right-skewed pattern, consistent with the strong positive correlation between distance and duration observed in the data.

Statistical Rationale for Log-Transformation

When a variable exhibits right skew (skewness > 1), the mean is pulled above the median by extreme values. In this dataset, a small number of very long-haul routes disproportionately inflate the arithmetic mean distance, leading to systematic ETA underestimation for typical short-haul shipments. Log-transformation compresses the scale of extreme values, aligns the mean and median, and produces a distribution that satisfies the approximate normality assumptions required by linear and kernel-based algorithms.

5.3.2 Correlation Heatmap Analysis

A **Pearson Correlation Heatmap** was constructed across all key numeric features — including transportation distance, log-distance, trip duration, average speed, delay days, the binary delay target, pickup hour, day of week,

weekend flag, and booking month. The upper triangle of the symmetric matrix was masked to reduce visual redundancy. The analysis revealed several operationally significant relationships:

- **Distance–Duration Correlation:** A strong positive correlation was observed between raw distance and trip duration, confirming that the SLA (Planned ETA) framework does account for route length, but the residual scatter around the regression line indicates that SLA buffers are insufficiently calibrated for long-haul routes.
- **Speed–Delay Inverse Relationship:** Average speed exhibited a meaningful negative correlation with the delay indicator: slower average speeds were associated with higher delay probability, validating speed as a potential real-time early-warning signal for at-risk shipments in the field.
- **Temporal Features:** Pickup hour and day-of-week showed modest but non-zero correlations with delay, confirming that departure timing carries genuine predictive signal beyond mere route characteristics.
- **Low Cross-Collinearity Among Engineered Features:** The log-distance and raw distance were, as expected, highly correlated with each other but not with temporal features, validating the feature set's information diversity and reducing concerns about multicollinearity in the classification model.

5.3.3 Delay Profiling by Vehicle Type and Temporal Dimension

A critical dimension of the exploratory analysis was the **profiling of delay rates across categorical operational strata**. For each vehicle type with at least 30 shipments, the mean **target_delay** rate was computed and visualised as a colour-coded bar chart (red for delay rates above 70%, amber for 50–70%, green below 50%). This analysis revealed substantial heterogeneity across vehicle categories: certain heavy commercial vehicle types exhibited delay rates exceeding 70%, while other vehicle classes maintained rates below 50% — a differential that cannot be explained by route length alone and points to vehicle-specific maintenance, GPS coverage, or operational management issues.

Temporal profiling further revealed that **late-night and early-morning bookings** carried elevated delay rates compared to daytime operations, likely reflecting reduced dispatcher oversight, driver fatigue effects, and lower road-assistance availability during off-hours. Weekend bookings showed a similarly elevated risk profile relative to weekday dispatches. These temporal patterns provide actionable intelligence for scheduling and workforce deployment decisions at the operational level.

5.4 Frequent Pattern Mining: The Apriori Algorithm

Univariate and bivariate analyses identify how individual variables relate to delay. However, logistics delay is rarely caused by a single factor in isolation — it is typically the product of a **compound combination** of circumstances: a specific vehicle class, operating on a long-haul corridor, with a particular GPS provider, departing late at night. To discover these hidden multi-factor delay patterns, the team applied the **Apriori Algorithm** — the foundational technique of market-basket analysis, adapted here for logistics transaction mining.

5.4.1 Algorithm Overview and Conceptual Foundation

The Apriori algorithm operates on a **transaction database** where each transaction (here, each shipment record) is represented as a set of items (here, categorical attribute-value pairs such as *VEH:Heavy_Commercial*, *GPS:Provider_A*, *STATUS:DELAYED*). The algorithm identifies **frequent itemsets** — combinations of items that co-occur in at least a minimum proportion of transactions (the *minimum support* threshold) — and then generates **association rules** of the form $A \rightarrow B$ from these itemsets. The strength of each rule is evaluated using three core metrics:

Metric	Mathematical Definition	Operational Interpretation
Support	$P(A \cap B)$	How frequently does this combination appear in the dataset? High support = pattern is not a rare coincidence.
Confidence	$P(B A)$	When condition A is present, how often does B also occur? High confidence = strong predictive reliability.
Lift	$\text{Confidence} / P(B)$	Is this combination more frequent than random chance? $\text{Lift} > 1.0$ = genuine positive association; the higher, the stronger.

5.4.2 Implementation and Transaction Construction

Each shipment record was encoded as a transaction basket containing items from four operational dimensions: **vehicle category** (prefixed *VEH:*), **material category** (prefixed *MAT:*), **GPS provider brand** (prefixed *GPS:*), and **delay status** (*STATUS:DELAYED* or *STATUS:ONTIME*). This prefix convention ensures that items from different attribute domains remain distinguishable, preventing spurious cross-domain associations. The **TransactionEncoder** from the *mlxtend* library was used to convert the item-list representation into the binary matrix format required by the Apriori implementation.

The Apriori algorithm was executed with a **minimum support threshold of 3%** (i.e., a combination must appear in at least 3% of all shipments to qualify as a frequent itemset) and a **minimum lift threshold of 1.05** for rule generation. The resulting rule set was then filtered to retain only rules where the *consequent* was *STATUS:DELAYED*, isolating rules with direct predictive relevance to the project's primary operational concern.

5.4.3 Key Association Rule Findings

The filtered association rules targeting **STATUS:DELAYED** as the consequent revealed a set of compound risk profiles with lift values substantially above 1.0, confirming that these combinations are not accidental co-occurrences. The most operationally significant findings include:

- GPS Provider–Vehicle Type Interactions:** Certain combinations of GPS hardware provider and vehicle category appeared with disproportionately high delay association. This pattern suggests that specific GPS

device families may suffer from incompatible firmware or poor update frequency on certain vehicle classes, leading to missed route deviations and uncorrected delays.

- **Long-Haul Routes with Heavy Vehicles:** Rules containing both a 'long-haul' distance bin indicator and a heavy commercial vehicle category consistently appeared at high lift values, indicating that scheduled SLA buffers systematically underestimate the elapsed time for this operational combination.
- **Specific Origin-Hub Patterns:** Certain source hub prefixes appeared repeatedly in delay-predicting rules, identifying origin hubs with structurally elevated risk — whether due to loading bottlenecks, poor road connectivity from the hub perimeter, or vehicle maintenance quality at those locations.
- **Material-Vehicle Mismatches:** Rules involving material categories typically transported by a mismatched vehicle type (e.g., temperature-sensitive goods on standard dry-freight vehicles) showed elevated delay lift, pointing to potential SLA violation risks from resource allocation inefficiencies.

Business Value of Association Rule Mining: Unlike correlation analysis, which measures the linear association between two variables in isolation, association rules capture the **joint effect of multiple simultaneous conditions**. A logistics dispatcher informed by a high-lift delay rule (e.g., *IF heavy vehicle type AND eastern corridor AND late-night departure THEN $P(\text{delay}) = 82\%$*) can take targeted pre-emptive action — reassigning the vehicle, adjusting the departure time, or setting a revised customer expectation — **before the shipment departs**. This constitutes a direct operational upgrade from reactive incident management to proactive risk governance.

5.5 Predictive Modelling: Random Forest Delay Classifier

The analytical capstone of the outbound system is a supervised machine learning classifier trained to predict the **target_delay** binary variable — indicating whether a given shipment will breach its planned SLA — using exclusively information available at the moment of booking. This design constraint, known in machine learning practice as the **no-future-leakage principle**, ensures that the model's outputs are operationally deployable: a dispatcher can score any new shipment at booking time and receive a risk assessment before the vehicle departs the warehouse.

5.5.1 Model Selection Rationale

Three classification algorithms were evaluated: **Logistic Regression** (linear baseline), **Random Forest Classifier** (ensemble of decision trees), and **Gradient Boosting Classifier** (sequential boosting ensemble). The Random Forest was selected as the primary reporting model on the basis of its balance of predictive performance, interpretability through feature importance scores, and robustness to overfitting. The Random Forest's ensemble architecture — aggregating predictions across 100 independent decision trees, each trained on a random bootstrap sample of the training data — makes it inherently resistant to the high-variance behaviour that afflicts single decision trees on noisy operational data.

A critical modelling consideration was the observed **class imbalance** in the dataset: the majority of shipments in this Indian logistics network were delayed, reflecting a systemic operational challenge rather than an isolated anomaly. To address this imbalance, the `class_weight='balanced'` parameter was applied to all classifiers, causing

the algorithm to assign proportionally higher penalty weights to misclassifying minority-class (on-time) instances. **Precision-Recall AUC (PR-AUC)** was adopted as the primary evaluation metric in preference to ROC-AUC, since PR-AUC is more informative under class imbalance conditions, as it is not inflated by the large number of true negatives in imbalanced binary problems.

5.5.2 Feature Engineering and Preprocessing Pipeline

A rigorous **scikit-learn Pipeline** architecture was constructed to ensure reproducible, leak-free data preprocessing. The pipeline comprised two parallel transformation tracks unified by a `ColumnTransformer`:

- **Numerical Track:** Features including `log_distance`, `pickup_hour`, `day_of_week`, `is_weekend`, and `booking_month` were processed through median imputation (robust to remaining outliers) followed by `StandardScaler` normalisation.
- **Categorical Track:** Features including vehicle category (`veh_cat`), origin state, destination state, and GPS brand were processed through constant-value imputation (filling missing categories with a designated 'missing' token) followed by `OneHotEncoder` transformation, producing a sparse binary representation suitable for tree-based models.

The dataset was partitioned into training (75%) and test (25%) sets using **stratified sampling** to preserve the class balance ratio across both partitions. Model evaluation was reinforced through **5-fold Stratified Cross-Validation**, which provides a statistically robust estimate of generalisation performance that is not inflated by a single favourable train-test split.

5.5.3 Model Performance and Interpretation

Post-training evaluation on the held-out test set confirmed that the **Random Forest Classifier** delivered meaningful predictive performance across all reported metrics. Confusion matrix analysis revealed the model's ability to correctly identify the majority of delayed shipments (high recall), with precision performance reflecting the inherent difficulty of distinguishing delayed from on-time shipments in the presence of noise from incomplete GPS data and external weather events not captured in the feature set.

The model's **feature importance analysis** — computed as the mean decrease in Gini impurity across all trees — identified the following as the most influential predictors of delay: log-transformed route distance (confirming the long-haul effect), vehicle category, destination state (reflecting regional infrastructure quality differentials), origin state, and GPS provider brand (corroborating the association rule finding that GPS hardware quality correlates with operational delay patterns).

Model	Primary Metric	Design Choice	Operational Role
Logistic Regression	Baseline PR-AUC	class_weight='balanced', liblinear solver	Linear interpretability benchmark
Random Forest (Primary)	Cross-validated PR-AUC	100 trees, balanced class weights, OOB scoring	Primary production scorer + feature ranking
Gradient Boosting	Best raw PR-AUC	100 estimators, lr=0.1, sequential residual fitting	High-performance benchmark for comparison

5.5.4 Operational Impact: From Reactive to Proactive Logistics

Strategic Shift Enabled by the Random Forest Model

Traditional logistics operations are inherently reactive: a shipment is flagged as delayed only after its planned ETA has been missed, at which point customer communication, re-routing, or compensation procedures must be initiated under time pressure and with limited options. The trained Random Forest Classifier fundamentally changes this dynamic. By scoring each new booking at the moment of dispatch planning — before the vehicle leaves the warehouse — the model produces a **delay risk probability** for every shipment. High-risk bookings (those exceeding a configurable probability threshold) can be surfaced immediately to dispatch supervisors, enabling interventions such as: vehicle reassignment, pre-emptive customer notification, route replanning, or adjusted SLA commitments — all before departure. This transforms the logistics operation from a *cost centre that manages failures* into a **risk-intelligent system that prevents them**.

CHAPTER 6: Visual Data Mining & Operational Bottlenecks Analysis

Supplementary Exploratory Analysis Using Orange Data Mining Software — Outbound Logistics Dataset

This chapter presents the findings of a supplementary visual data mining and exploratory analysis performed on the outbound logistics dataset using **Orange Data Mining software**. The objective of this investigation was to complement the quantitative modelling approaches documented in preceding chapters by leveraging the interpretive power of visual analytics to surface operational patterns, anomalies, and structural inefficiencies that may otherwise remain concealed within aggregate statistical summaries.

The analysis is organised around four critical operational dimensions: **(1)** outlier detection through trip-distance-to-duration scatter analysis; **(2)** temporal demand characterisation across the working week; **(3)** vendor-specific delay profiling to differentiate systemic failure from statistical noise; and **(4)** an ETA feasibility assessment correlating route distance with observed delay rates. The combined findings yield actionable business intelligence with direct implications for **fleet scheduling, vendor governance, resource allocation, and ETA model recalibration**.

6.1 Visual Outlier Detection — Trip Distance vs. Duration Analysis

6.1.1 Methodology

A scatter plot was constructed in Orange with **Trip Distance (km)** plotted on the horizontal axis (X) and **Trip Duration (hours)** plotted on the vertical axis (Y). Individual shipments were represented as discrete data points, while **vehicle type was encoded as a colour dimension** to enable simultaneous multi-variable visual

interpretation. This approach leverages the human visual system's capacity to detect patterns, clusters, and deviations that would be difficult to identify through tabular data review alone.

6.1.2 Core Finding — General Distribution of Trip Distances

The scatter plot revealed that the **overwhelming majority of outbound shipments cover distances of less than 500 km**. This concentration in the short-to-medium haul range is consistent with regional distribution network patterns. The dense cluster in the lower-left quadrant of the plot establishes the **operationally normal baseline** — a reference boundary against which deviant observations can be objectively assessed.




Figure 6.1: Scatter plot illustrating severe trip duration outliers despite relatively short transportation distances.

6.1.3 Critical Finding — Severe Outliers Indicating Hidden Operational Bottlenecks

Against this established baseline, the scatter plot exposed a subset of **severe operational outliers** of immediate investigative concern. Specifically, a number of shipments covering distances of approximately **500 km** were **recorded as consuming in excess of 350 hours** to complete. Under standard operating conditions, a 500 km route should be executable within a range of approximately 8 to 20 hours depending on vehicle type and road category. A recorded duration exceeding 350 hours represents a **delay multiplier of 17x to 44x** the expected norm — a deviation of sufficient magnitude to classify these events as operational crises rather than ordinary variance.

These extreme outliers are not attributable to distance alone. The visual evidence points strongly to one or more of the following **hidden operational bottlenecks**:

- **Severe Traffic Congestion:** Persistent gridlock or road infrastructure failure on specific corridors not captured within the existing ETA model inputs.
- **Vehicle Mechanical Breakdowns:** Extended unscheduled maintenance stops or complete in-transit vehicle failures, indicating gaps in preventive maintenance protocols and roadside contingency planning.
- **Critical Loading and Unloading Delays:** Prolonged dwell times at origin or destination facilities caused by dock capacity constraints, scheduling conflicts, or administrative processing bottlenecks.
- **Regulatory and Compliance Holds:** Unanticipated detentions at checkpoints, inspection facilities, or clearance points that artificially inflate total recorded trip duration.

Business Recommendation — Outlier Root-Cause Investigation

All shipments exhibiting a duration-to-distance ratio exceeding a defined anomaly threshold (e.g., >30 hours per 100 km) should be **immediately flagged and escalated** for structured root-cause investigation. A dedicated operations task force should audit the specific routes, vehicle assignments, driver records, and loading timestamps associated with these outlier events. Findings should directly inform revisions to preventive maintenance schedules, route planning algorithms, and origin-facility capacity management protocols.

6.1.4 Vehicle Type Analysis — Colour-Encoded Performance Segmentation

The vehicle type colour dimension embedded within the scatter plot provided an additional layer of diagnostic intelligence. As operationally expected, **heavy trailers** dominated the long-haul segments (routes exceeding 500 km), consistent with their design capacity for inter-city freight. However, the colour segmentation also served a comparative performance function: it enabled the identification of **specific vehicle types that are systemically slower than fleet peers** across equivalent distance bands.

Vehicle types whose data points consistently occupy the upper boundary of the duration distribution for a given distance range should be classified as **"High-Latency Fleet Assets"**. These vehicles warrant mechanical assessment, route reassignment, or retirement scheduling. This analysis thereby transforms a conventional scatter plot visualisation into a **data-driven fleet performance audit instrument**.

6.2 Temporal Demand Patterns — Day-of-the-Week Analysis

6.2.1 Methodology

A bar chart was generated to examine the total distribution of outbound shipment orders across the days of the operational week. Each bar represents the aggregate order volume attributed to a specific day number, enabling direct visual comparison of daily operational load. This temporal analysis constitutes a prerequisite input for informed **fleet deployment planning, staff scheduling, and capacity optimisation**.

6.2.2 Peak Performance Day — Day 1 (Monday)

Day 1 of the operational week recorded the highest single-day order volume, with over 600 outbound shipments. This Monday surge is consistent with well-documented logistics demand patterns driven by the accumulation of weekend orders, the resumption of B2B procurement cycles, and the weekly restocking rhythms of downstream customers and retail partners.

The operational implications are multi-dimensional and require coordinated advance planning across fleet, warehouse, and customer service functions:

- **Maximum Fleet Pre-Positioning:** The full operational vehicle fleet should be serviced, fuelled, and physically pre-positioned at dispatch facilities by Sunday evening to ensure zero-delay readiness for Monday morning dispatch windows.
- **Full Staffing Activation:** Warehouse picking, packing, loading, and dispatch teams must be scheduled at maximum headcount on Day 1, with pre-approved overtime provisions available for activation if actual volumes exceed forecast.
- **Priority Dock Scheduling:** All loading dock allocations should be pre-scheduled and optimised for Monday to eliminate vehicle queuing and minimise loading dwell times during peak hours.

- **Proactive Customer Communication:** Customers submitting orders on Day 1 should receive automated notifications acknowledging receipt and setting realistic processing timelines, managing expectations transparently during the high-demand window.




Figure 6.2: Weekly demand distribution highlighting peak operations on Day 1 and a significant drop on Day 6.

6.2.3 The Weekend Gap — Day 6

In sharp contrast to the Monday peak, **Day 6 recorded a dramatic reduction in operational volume**, with shipment orders declining to approximately 120 — representing less than 20% of Monday's throughput. This Weekend Gap is a structural and recurring feature of the demand cycle that, if managed proactively, presents a strategic opportunity rather than an inherent operational inefficiency.

- **Scheduled Preventive Fleet Maintenance:** With minimal operational pressure on Day 6, the maintenance department should execute comprehensive vehicle inspections, scheduled servicing, and any outstanding minor repairs. This practice directly addresses the breakdown-related outliers identified in Section 8.1 and ensures peak-day fleet reliability.
- **Capacity Right-Sizing and Personnel Cost Reduction:** Staffing levels on Day 6 should be formally reduced through structured part-time scheduling or a rotating reduced-headcount roster, generating measurable personnel cost savings while maintaining minimum coverage obligations.

6.2.4 Mid-Week Operational Stability — Days 0, 2, 3, and 4

Days 0, 2, 3, and 4 demonstrated a high degree of operational consistency, with order volumes maintaining a stable band of approximately **550 to 600 shipments per day**. This mid-week stability is a significant positive indicator: it confirms that the logistics system operates within a **predictable, controllable operational rhythm** during the bulk of the working week, thereby enabling reliable long-range planning, supplier synchronisation, and resource pre-commitment.

Weekly Demand Profile — Strategic Resource Allocation Summary

The following demand pattern should govern all weekly resource planning decisions:

- **Day 1 (Monday):** Peak day — >600 orders. Activate full fleet and maximum staffing.
- **Days 0, 2, 3, 4:** Stable period — 550–600 orders. Standard operations apply.
- **Day 6 (Weekend):** Low-volume day — ~120 orders. Schedule fleet maintenance and implement staffing reductions.

6.3 Vendor-Specific Delay Profiling — Systemic Failure vs. Statistical Noise

6.3.1 Methodology and Analytical Framework

A vendor-level delay analysis was conducted by cross-tabulating each company's total recorded shipment volume against their respective late delivery rate. The resulting visualisation identified multiple companies exhibiting a **100% delay rate**, prompting a second-order analytical question of critical business importance: do these rates reflect genuine operational dysfunction, or are they merely statistical artefacts produced by insufficient sample sizes? Conflating these two categories would risk misallocating management resources and damaging productive vendor relationships on the basis of inadequate evidence.

6.3.2 Category 1 — Statistical Noise (Low-Sample Artefacts)

The analysis identified two vendors displaying a 100% delay rate that must be interpreted with significant caution due to their critically small sample sizes:

Company Name	Total Shipments	Recorded Delay Rate	Classification
amalgamations valeo clutch pvt	1	100%	Insufficient Data — Statistical Noise
jbm ma automotive pvt ltd	2	100%	Insufficient Data — Statistical Noise

A 100% delay rate derived from one or two shipment observations carries **no statistically valid inferential weight**. These figures should not trigger vendor sanctions, contract termination procedures, or formal escalation protocols. The appropriate response is to **expand the observation window** and accumulate a meaningful sample volume before drawing performance conclusions. Premature punitive action risks damaging potentially productive and reliable vendor partnerships on the basis of anecdotal evidence.

6.3.3 Category 2 — Systemic Failure (High-Volume, Immediately Actionable)

In direct contrast, two additional vendors exhibiting high delay rates were supported by statistically meaningful shipment volumes, elevating their delay rates from probabilistic artefacts to **confirmed, actionable systemic failures** demanding executive-level intervention:

Company Name	Total Shipments	Recorded Rate	Delay	Risk Classification		
Otis Elevator Company (India) Ltd	34	100%	(34 / 34 delayed)	CRITICAL	—	Confirmed Systemic Failure
Wipro	63	98.4%	(62 / 63 delayed)	CRITICAL	—	Confirmed Systemic Failure

Otis Elevator Company (India) Ltd accumulated 34 recorded shipments, **every single one of which was delivered late**. The probability of 34 consecutive delays occurring by chance, in the absence of a structural causal factor, is statistically negligible. The evidence points unambiguously to a systemic issue embedded within either the specific routes assigned to Otis, the loading and scheduling procedures at their facilities, or the contractual SLA frameworks governing their delivery commitments.

Wipro accumulated 63 shipments — one of the highest vendor volumes in the dataset — of which **62 were recorded as delayed**, yielding a 98.4% delay rate. At this volume, the cumulative operational cost of these delays — encompassing penalty clause exposure, customer satisfaction deterioration, and downstream supply chain disruption — is substantial and demands immediate commercial remediation.

⚠️ Critical Vendor Governance Recommendation — Otis Elevator Company & Wipro

The following structured interventions are recommended for immediate execution: **1. Emergency Contract and SLA Audit:** Both vendors' logistics agreements should be reviewed immediately to assess whether current delivery windows reflect operationally achievable commitments given observed performance data. **2. Root-Cause Joint Working Sessions:** Structured problem-solving workshops should be convened with each vendor's operations team to map the end-to-end shipment process and isolate specific delay-generating stages (e.g., loading, in-transit, unloading, administrative). **3. Route-Level Disaggregation:** Delay patterns should be segmented by route corridor, time of day, and vehicle type to determine whether the failure is geography-specific or process-specific in nature. **4. Formal Performance Improvement Plans (PIPs):** Both vendors should receive a formal PIP document specifying measurable KPI targets, review intervals, and clearly defined escalation consequences — including contract termination provisions — if benchmark improvements are not achieved within the stipulated timeframe.

6.3.4 The Zero-Delay Benchmark

The analysis additionally surfaced a cohort of vendors maintaining a **0% delay rate**. These companies represent operational best-in-class performance within the current network and constitute a valuable internal benchmarking resource. A structured study of their logistical practices, route configurations, facility layouts, and contractual terms — consistent with a **positive deviance methodology** — could yield transferable insights applicable across the broader vendor portfolio for systematic performance uplift.



Figure 6.4: Correlation between trip distance and delay rates, indicating a flawed ETA model for long-haul shipments.

6.4.3 Secondary Anomaly — Short-Distance, High-Duration Outliers

The scatter plot simultaneously exposed a distinct and analytically important sub-pattern: a cluster of data points appearing at relatively short distances but exhibiting disproportionately extended durations, in some cases exceeding **250 hours**. These outliers are geographically inexplicable — the route distance alone provides no plausible justification for the observed duration, eliminating infrastructure constraints as a causal explanation.

These short-distance, high-duration events signal an **operational efficiency failure** attributable to specific companies, individual drivers, or localised process breakdowns — distinct in nature from the ETA model inadequacy identified for long-haul routes. They represent a **reputational and financial liability disproportionate to their geographic complexity** and must be treated as a separate, dedicated investigation category.

Strategic Recommendation — ETA Model Recalibration

1. Immediate Long-Haul ETA Revision: All routes exceeding 1,000 km should have their ETA calculations reviewed and recalibrated using actual historical trip duration data. A distance-band model is recommended, applying differentiated ETA multipliers for routes in the 500–1,000 km, 1,000–1,500 km, and >1,500 km ranges. **2. Short-Distance Outlier Audit:** A dedicated extract of all shipments with distances under 500 km but durations exceeding 250 hours should be generated and assigned to the operations team for vendor and driver accountability investigations. These cases require personnel performance measures, not algorithmic adjustments. **3. Customer Expectation Management:** Pending full ETA model recalibration, all customer-facing delivery windows for routes exceeding 1,000 km should be conservatively extended to reflect realistic completion times. This measure will reduce formal delay counts, improve customer satisfaction metrics, and reduce commercial penalty clause exposure during the recalibration transition period.

6.5 Chapter Summary — Consolidated Action Intelligence Matrix

The table below consolidates the four visual data mining findings into a unified action intelligence matrix, mapping each analytical insight to its identified root cause, risk classification, and corresponding recommended corrective action.

Section	Analytical Finding	Risk Level	Recommended Action
6.1	Outlier trips: ~500 km routes exceeding 350 hours duration	● Critical	Root-cause audit of specific routes, vehicles, loading procedures, and driver logs
6.1	Heavy trailers systemically slower on long-haul routes	● Moderate	Fleet performance segmentation; schedule high-latency assets for mechanical review
6.2	Day 1 (Monday) peak — >600 outbound orders	● Planning	Pre-position fleet Sunday; activate maximum staffing and priority dock scheduling
6.2	Day 6 (Weekend) — volume drops to ~120 orders	● Opportunity	Schedule preventive maintenance; implement reduced staffing and cost right-sizing
6.3	Otis Elevator: 34 shipments, 100% delayed	● Critical	Emergency SLA review; issue formal Performance Improvement Plan (PIP)
6.3	Wipro: 63 shipments, 98.4% delayed	● Critical	Emergency SLA review; issue formal Performance Improvement Plan (PIP)
6.4	Delay rate increases proportionally with trip distance	● Critical	Recalibrate ETA model for all routes >1,000 km using historical duration data
6.4	Short-route trips recording >250-hour durations	● High	Dedicated driver and vendor efficiency audit; apply personnel accountability measures

In summary, the visual data mining exercise conducted through Orange Data Mining software has systematically transformed raw outbound logistics data into a coherent set of operational intelligence signals. These signals — spanning fleet performance anomalies, temporal demand structure, vendor accountability deficiencies, and ETA model validity — collectively define a **structured and prioritised improvement roadmap** for the logistics intelligence system. The findings presented in this chapter will directly inform the system design recommendations, KPI dashboard specifications, and strategic intervention frameworks developed in the subsequent chapters of this report.

Chapter 7: Strategic Business Recommendations

The dual analytical systems developed in this project — the Power BI-driven Warehouse Intelligence System (Inbound) and the Python/SQL-driven Delivery Tracking Analytics System (Outbound) — have collectively produced a rich, evidence-based picture of where operational inefficiencies reside, how they interact, and where targeted investment will generate the highest return. This chapter translates those analytical findings into a structured set of strategic recommendations addressed to the company's executive leadership and operations management teams.

Each recommendation is grounded in a specific quantitative finding from the project's analysis, articulates a clear mechanism of expected impact, and is framed as an implementable operational policy rather than a general

aspiration. The recommendations are ordered by strategic priority, from immediate operational fixes to longer-term technology integration.

Recommendation 1: Integrate the Machine Learning Risk Scorer into the Daily Dispatch Workflow

Priority Level: CRITICAL | Implementation Timeline: 0–3 Months

Data-Driven Foundation: The Random Forest Classifier, validated through 5-fold stratified cross-validation, is capable of assigning a delay risk probability to every new shipment at the point of booking — using only information available before departure.

The Problem: Under current operations, delay risk is identified only after an SLA breach has occurred, forcing the organisation into costly reactive mode: expedited deliveries, compensatory measures, and damage to customer relationship metrics. Every delayed shipment that could have been predicted in advance represents a preventable failure.

The Recommendation: The trained and serialised Random Forest Pipeline (saved as a *.pkl* artefact using *joblib*) should be deployed as a **Booking-Time Risk API** integrated into the company's Transport Management System (TMS). Each new shipment booking should trigger an API call that returns a **delay probability score (0–1)** and a **risk tier (Low / Medium / High)** based on a configurable threshold (recommended starting threshold: 0.65).

Operational Implementation Steps:

- 1. Export the trained sklearn Pipeline object and wrap it in a lightweight REST API endpoint (FastAPI or Flask) callable from the existing TMS booking interface.
- 2. Define risk tiers: Low Risk (< 0.40), Medium Risk (0.40–0.65), High Risk (> 0.65). Display these on the dispatcher dashboard in colour-coded format.
- 3. For High-Risk bookings, automatically trigger a supervisor alert with the top three contributing risk factors (derived from SHAP feature importance analysis) to enable targeted intervention.
- 4. Log all model predictions and actual outcomes to create a continuously growing labelled dataset for quarterly model retraining and performance monitoring.

Expected Business Impact: A conservative assumption of reducing high-risk shipment delays by even 20% would, at scale, translate into measurable improvements in on-time delivery KPIs, reduced customer escalation volumes, and lower expediting costs — outcomes directly visible to executive stakeholders within a single fiscal quarter.

Recommendation 2: Reconfigure Vehicle-to-Route Assignments Based on Delay Profiling

Priority Level: HIGH | Implementation Timeline: 1–3 Months

Data-Driven Foundation: EDA identified that certain vehicle categories exhibit persistently elevated delay rates (exceeding 70%) independent of route length, while other vehicle classes operating comparable distances maintain delay rates below 50%. The Apriori analysis further confirmed that specific vehicle type + route corridor combinations are disproportionately represented in high-lift delay association rules.

The Problem: Current vehicle allocation practices appear to assign vehicles to routes based primarily on capacity and availability rather than on the vehicle type's historical performance profile on specific route types. This leads to a predictable but unmanaged pattern of vehicle-route mismatches that systematically elevate delay rates on certain corridors.

The Recommendation: The operations team should construct a **Vehicle–Route Performance Matrix** — a lookup table, directly derivable from the project's EDA and association rule outputs — that maps each vehicle category to its historical delay rate on each route corridor type (short-haul, medium-haul, long-haul inter-state). This matrix should be used to inform vehicle assignment logic in the TMS, with the following specific policies:

- **Route-Specific Vehicle Reservation:** Vehicle types with delay rates exceeding 70% on long-haul inter-state routes should be restricted from those assignments and reallocated to short urban-delivery circuits where their delay profiles are more acceptable.
- **Weather-Condition Vehicle Policy:** Association rules involving weather-sensitive route conditions (monsoon-affected corridors in particular) and heavy vehicle types with known delay associations should trigger automatic vehicle class upgrades or additional SLA buffer allocation at the time of booking. Lighter, more manoeuvrable vehicle categories demonstrated lower delay rates on weather-affected routes in the EDA analysis.
- **Supplier-Level Vehicle Standards:** Supplier performance analysis revealed significant inter-supplier variation in on-time delivery rates. Suppliers whose vehicle fleets show persistently high delay profiles on specific routes should be subject to a structured performance improvement programme, with specific vehicle class standards stipulated in revised service contracts.

Recommendation 3: Redesign Warehouse Zone Layouts Based on Picking Time and Velocity Data

Priority Level: HIGH | Implementation Timeline: 2–4 Months

Data-Driven Foundation: The Inbound Power BI analysis identified statistically significant variation in picking times across warehouse storage zones. High-velocity SKUs (items with the highest outbound dispatch frequency) were found to be disproportionately concentrated in suboptimal storage zones — those requiring longer picker travel distances — causing unnecessary loading delays that cascade into delayed vehicle departures.

The Problem: Warehouse zone layout decisions are often made at facility setup and rarely revisited as inventory composition and demand patterns evolve. When high-velocity SKUs are stored far from dispatch docks, pickers accumulate unnecessary travel time, depressing their effective throughput rate and creating a queuing bottleneck that delays vehicle loading — particularly during peak dispatch windows.

The Recommendation: A formal **Velocity-Based Slotting Optimisation** exercise should be conducted using the Power BI picking time and inventory turnover data as its input. The guiding principle of velocity-based slotting is straightforward: **the highest-frequency items should be stored in the locations closest to the dispatch dock**, minimising average picker travel distance and maximising throughput during peak loading periods. The following implementation framework is recommended:

- **ABC Inventory Classification:** Classify all SKUs into velocity tiers — A (top 20% by dispatch frequency, accounting for ~80% of movements), B (next 30%), and C (remaining 50%). This classification should be refreshed monthly using Power BI's automated data pipeline.
- **Zone Reallocation:** A-tier SKUs should be co-located in Zone 1 (the zone adjacent to the primary loading dock, identified in the Power BI zone analysis as having the fastest existing picking times). B-tier SKUs occupy Zone 2, with C-tier SKUs assigned to the furthest zone where extended picking times are acceptable given their low movement frequency.
- **Stockout Prevention Integration:** The Power BI safety stock and reorder point analysis identified the SKUs most frequently experiencing stockout conditions. These items should be given priority zone assignment within their velocity tier, as stockouts compound the picking-time delay problem by forcing pickers to seek substitute items or flag unavailability — a double penalty on outbound throughput.
- **Continuous Monitoring:** A Power BI KPI dashboard should track Average Picking Time per Zone on a weekly basis. Any zone with a picking time trend deteriorating beyond a defined control limit should automatically generate an alert for warehouse managers to review slotting assignments.

Recommendation 4: Replace Underperforming GPS Providers on High-Risk Routes

Priority Level: MEDIUM-HIGH | Implementation Timeline: 3–6 Months

Data-Driven Foundation: The Apriori association rule analysis identified specific GPS provider brands as components of high-lift delay prediction rules. The correlation heatmap and vehicle-type delay profiling further confirmed that GPS data quality — measured by update frequency and coordinate accuracy — correlates with delay detection latency and operational responsiveness.

The Problem: GPS telematics data is not merely a tracking amenity — it is the operational foundation of real-time logistics management. When GPS providers deliver infrequent updates, poor coordinate accuracy, or device failures, dispatchers lose visibility into shipment status, cannot detect route deviations in time to correct them, and lack the data signals that the predictive model relies upon to generate accurate delay risk scores. In effect, poor GPS hardware degrades both the operational responsiveness of the dispatch team and the predictive accuracy of the ML system.

The Recommendation: The company should conduct a **GPS Provider Performance Audit** — a structured evaluation of each telematics provider's contribution to delay outcomes — using the association rule and EDA findings from this project as a baseline. The audit should quantify:

- **Delay Rate Differential:** The difference in average delay rate between shipments equipped with each GPS provider versus the network average, controlling for route type and vehicle class.
- **Data Gap Frequency:** The proportion of shipments with each provider that experienced gaps in GPS pings exceeding 60 minutes, which would indicate device failure or connectivity loss during transit.
- **Anomaly Detection Cross-Reference:** The proportion of Isolation Forest anomaly flags that occur on shipments equipped with each GPS provider — anomalous speed and duration profiles are often caused by GPS data failures rather than genuine journey irregularities.

Providers that score adversely on two or more of these dimensions should be subject to a **time-bound performance improvement requirement**, with migration to alternative providers initiated for high-risk route corridors. The estimated cost of GPS hardware replacement on the most critical corridors should be evaluated against the customer satisfaction and SLA penalty cost savings that improved tracking quality would enable.

Recommendation 5: Establish a Unified Supply Chain Intelligence Dashboard and Governance Framework

Priority Level: MEDIUM | Implementation Timeline: 4–8 Months

Data-Driven Foundation: The current project demonstrated that the inbound warehouse system (Power BI) and the outbound delivery system (Python/SQL ML) produce complementary but currently siloed insights. The project's Dual-System architecture reveals a strategic gap: warehouse loading delays caused by poor slotting (an inbound problem) contribute to vehicle departure delays that cascade into delivery SLA breaches (an outbound problem) — a causal chain that neither system currently captures end-to-end.

The Problem: When inbound and outbound analytics exist in separate systems with no shared data model, root-cause analysis becomes fragmented: the outbound team sees a delayed shipment but cannot easily attribute it to a warehouse loading bottleneck; the warehouse team has no visibility into how their picking time performance affects downstream customer delivery metrics. This organisational and technological silo is a structural inhibitor of supply chain optimisation.

The Recommendation: The organisation should invest in a **Unified Supply Chain Intelligence Platform** that integrates both the warehouse operations data (currently managed in Power BI) and the outbound delivery analytics data (currently managed in Python/SQL) into a single shared data warehouse and executive dashboard. The recommended architecture consists of four layers:

- **Data Integration Layer:** A lightweight ETL pipeline that consolidates picking time records, zone inventory levels, vehicle departure times, and GPS tracking events into a single star-schema data warehouse (e.g., PostgreSQL or cloud-hosted BigQuery/Snowflake), updated daily.
- **Predictive Scoring Layer:** The deployed Random Forest Pipeline integrated as a scheduled batch scoring job, enriching each new booking record with a delay probability score before the morning dispatch planning session.
- **Executive KPI Dashboard Layer:** A Power BI dashboard consuming the integrated warehouse exposing end-to-end KPIs including: Warehouse-to-Vehicle Loading Time, Departure Punctuality Rate, On-Time Delivery Rate by Route and Vehicle Type, and Predicted vs Actual Delay Rate — providing executives with a single coherent view of supply chain health from storage to doorstep.
- **Governance and Continuous Learning Layer:** A quarterly model retraining cadence using accumulated prediction logs, a data quality monitoring process that alerts data engineering teams to anomalous GPS data gaps, and a formal model versioning process that allows the organisation to track and audit changes in model performance over time.

Long-Term Strategic Value: The Unified Intelligence Platform described above represents more than a technology upgrade — it is the institutional infrastructure for a **data-driven logistics culture**. Organisations that systematically close the loop between predictive analytics and operational decision-making are demonstrably better positioned to respond to supply chain disruptions, negotiate from strength in carrier and supplier contracts, and deliver the customer experience reliability that drives long-term commercial advantage in the logistics sector.

Chapter Summary

This chapter has translated the analytical findings of the End-to-End Transportation & Logistics Tracking Intelligence System into five concrete, prioritised, and evidence-grounded strategic recommendations for the executive leadership team. Taken together, these recommendations describe a coherent transformation roadmap that moves the organisation from its current state — characterised by reactive delay management, siloed analytical systems, and under-utilised GPS and operational data — to a future state defined by proactive ML-driven risk governance, optimised warehouse and vehicle resource allocation, and an integrated intelligence platform that generates compounding operational value over time.

The dual-system architecture of this project — combining Power BI warehouse analytics with Python-driven delivery prediction — has demonstrated that the most valuable logistics insights emerge not from any single technique in isolation, but from the deliberate integration of descriptive, diagnostic, and predictive analytics across the full order-to-delivery supply chain. The implementation of these recommendations represents the logical and strategic next phase of that integration journey.

CHAPTER 8: Conclusion

This project has demonstrated, in rigorous technical and operational detail, that the transformation of a logistics enterprise from a reactive, experience-driven organisation into a genuinely intelligence-led operation is not a theoretical aspiration — it is a structurally achievable outcome, provided the right architectural foundation is in place. Beginning with the disciplined engineering of the **LogisticT** SQL Server relational database — a rigorously normalised, constraint-enforced, eleven-table schema designed to Third Normal Form and enriched with a suite of ten Business Intelligence SQL VIEWS — the project established an unimpeachable data governance foundation upon which every subsequent analytical layer was built. From this bedrock, the inbound analytical dimension was addressed through a **Power BI** enterprise intelligence dashboard that exposed a financially material Order Fulfillment Rate deficit of 10.03 percentage points against the industry benchmark, identified 237 critical SKUs operating below their designated reorder thresholds, and benchmarked operational performance across four warehouse zones — revealing Zone A as the internal standard of excellence and quantifying the cascading consequences of sub-benchmark performance in Zones B, C, and D. The outbound analytical dimension was addressed through a Python-driven predictive pipeline that combined Haversine-based geospatial imputation, advanced Exploratory Data Analysis, Apriori association rule mining, and a validated Random Forest Classifier capable of assigning delay risk probabilities to new shipments at the moment of booking. The supplementary visual data mining investigation conducted in **Orange** completed the analytical arc by surfacing fleet performance anomalies, temporal demand rhythms, vendor accountability failures at the systemic level, and the structural inadequacy of the existing ETA model for long-haul routes exceeding 1,000 kilometres. Taken together, these four interconnected analytical layers constitute a complete, production-ready intelligence system spanning the full order-to-delivery lifecycle of a modern logistics operation.

The business value delivered by this **Dual-System Architecture** is best understood not as the sum of its individual components, but as the product of their deliberate integration. Each of the five strategic business problems defined in Chapter 2 has been directly addressed by a corresponding analytical output. The absence of a predictive risk management capability — the most consequential operational gap identified at the outset — has been resolved by a trained and serialisable machine learning pipeline that transforms booking-time data into actionable delay probability scores, enabling dispatch supervisors to intervene before a vehicle departs the warehouse rather than managing failure after it has occurred. The systemic pattern of SLA breach attributable to a structurally flawed ETA model has been exposed and quantified, with a tiered recalibration framework recommended for all routes segmented by distance band. The near-total failure rates of vendors such as Otis Elevator Company and Wipro — representing confirmed systemic non-performance rather than statistical variance — have been surfaced with the evidentiary rigour required to initiate formal Performance Improvement Plans. The hidden financial exposure of the warehouse inventory network — estimated at approximately **\$38.67 million per month in revenue at risk** — has been rendered visible, attributable, and actionable through the Power BI financial impact assessment. And the previously invisible causal chain linking inbound warehouse inefficiency to outbound delivery failure has been formally established, measured, and made operationally exploitable. What the system delivers, in aggregate, is precisely what Chapter 2 called for: a **360-degree operational intelligence view** that transforms fragmented,

siloes, and underutilised operational data into a unified, evidence-based intelligence platform — one that serves the warehouse manager, the dispatch supervisor, the financial controller, and the executive leadership team with equal analytical rigour and equal operational relevance.

An important dimension of this project that merits explicit acknowledgement is the role played by advanced **Large Language Models (LLMs)** as analytical co-pilots throughout the development lifecycle. Specifically, **Google Gemini Pro** and **Anthropic Claude AI** were engaged at critical junctures across the project to accelerate and augment the team's technical capabilities. These AI assistants contributed meaningfully to the optimisation of complex multi-table SQL queries, the debugging and refinement of Python data engineering scripts, the validation of feature engineering logic, and the interpretation of nuanced statistical outputs in the context of their operational significance. Rather than replacing the analytical judgment of the project team, these LLMs functioned precisely as the concept of **Augmented Analytics** prescribes: as intelligent accelerators that compress iteration cycles, surface alternative methodological approaches, and elevate the overall technical quality of the final deliverables. The deliberate and transparent integration of these AI capabilities into the project workflow is not a concession of analytical authorship — it is a demonstration of the professional competence required to operate effectively at the frontier of modern data science, where the ability to leverage advanced AI tools in service of rigorous analytical outcomes is itself a core practitioner skill. This project thereby serves as a practical proof of concept for the transformative potential of human-AI collaboration in enterprise data engineering and predictive analytics workflows.

The competitive landscape of the global logistics industry is undergoing a structural reconfiguration. In an environment defined by compressed delivery windows, escalating customer expectations for real-time transparency, and increasing pressure on cost and SLA compliance, the organisations that will command sustainable market advantage are those that have transformed their operational data from a passive record of past events into an active, predictive instrument of future decision-making. The End-to-End Transportation & Logistics Tracking Intelligence System documented in this report provides precisely that capability. By integrating a governed relational data infrastructure, a descriptive warehouse intelligence layer, a predictive machine learning scoring system, and a visual data mining investigative framework into a single coherent analytical platform, this system equips logistics organisations with the foundational capability to anticipate risk before it materialises, allocate resources with statistical precision, and engage customers with the proactive transparency that converts service reliability into long-term commercial loyalty. As the volume and velocity of logistics data continue to expand — driven by the proliferation of IoT telematics, GPS tracking density, and real-time demand signals — the architectural approach demonstrated in this project is designed to scale alongside that growth, accommodating richer feature sets, more granular predictive models, and deeper integration across the full supply chain network. The implementation of an intelligence system of this scope and rigour is no longer a strategic luxury reserved for the largest global logistics enterprises. It is the operational baseline that separates the organisations that lead their markets from those that merely follow them — and this project has demonstrated, in concrete and measurable terms, precisely how that baseline is built.

Appendix B

A) Dataset Overview & Types

Q1) Dataset Overview

- **Shape:** 6,880 Rows, 32 Columns.
- **Key Columns Include:** BookingID, BookingID_Date, vehicle_no, Origin_Location, Destination_Location, Planned_ETA, actual_eta, ontime, TRANSPORTATION_DISTANCE_IN_KM, customerID, Material_Shipped.

Q2) Column Understanding

- **BookingID:** Unique identifier for the shipment (Categorical) - Must be unique.
- **Origin/Destination_Location:** Geographical addresses (Categorical).
- **BookingID_Date:** Order creation date (Datetime).
- **trip_end_date:** Delivery completion date (Datetime).
- **ontime:** Delivery Status flag (0 = Delayed, 1 = On Time) (Categorical).
- **TRANSPORTATION_DISTANCE_IN_KM:** Indicator of shipping distance and cost (Numeric) - Must be positive.

Q3) Data Types Check & Fixes

- **Observation:** Several date columns (BookingID_Date, trip_end_date) were initially stored as 'object' (strings).
 - **Action Taken:** Converted to proper Date/Time types using `pd.to_datetime()` to enable time-series analysis.
-

B) Data Quality (Missing, Duplicates, Validity)

Q4 & Q5) Missing Values Overview & Strategy

- Status_Text: 100% missing (Requires derivation from other columns).
- ontime: 62.96% missing.
- delay: 36.88% missing.
- Driver_MobileNo: 60.88% missing.
- **Strategy:** High missing rates in critical status columns validate the necessity of the advanced data imputation and feature engineering pipelines we developed and applied in Chapter 5.

Q6 & Q7) Duplicates Check

- **Exact Duplicate Rows:** 0
- **Duplicate Booking IDs (Primary Key):** 5 (Action: Dropped duplicates keeping the latest entry to maintain database integrity).

Q8) Validity Rules



- **Negative Costs/Distances:** 0 (Data is clean in this regard).
- **Logical Errors:** 65 deliveries were recorded with an actual_eta before the BookingID_Date. These anomalies were flagged for removal during the cleaning phase.

C) Univariate EDA (Single Column Insights)

Q10) Numeric Summary (Transportation Distance in KM)

- **Mean:** 553.86 KM
- **Median (50%):** 160.00 KM
- **Standard Deviation:** 758.98 KM
- **Max:** 2,954.70 KM

Q11 & Q12) Histograms & Boxplots (Outlier Detection)




- **Interpretation:** Transportation distance shows a highly right-skewed distribution. Outliers exist for very long distances ($> 1,500$ KM). However, these are valid long-haul shipments (e.g., cross-state logistics) and were kept in the dataset as they represent critical operational scenarios.


D) Relationships (Multivariate / Bivariate)

Q15 & Q19) Numeric Relationships & Heatmap

- **Top Correlation:** There is a notable correlation (0.136) between TRANSPORTATION_DISTANCE_IN_KM and Minimum_kms_to_be_covered_in_a_day.



Q16) Visual Check: Scatter Plots



Interpretation: The scatter plot appears empty because the categorical variable used for the color hue ('Status_Text', derived from 'On_time_Delivery') is 100% missing in the raw dataset, as identified in Q4. It is impossible to visually analyze the relationship with a status that does not exist in the raw data. This is a critical and successful finding of the Data Quality analysis. It proves that a simple analysis is impossible and strongly validates the need for the advanced data imputation and feature engineering pipelines we developed and applied in the core chapters of this report to reconstruct these missing statuses.

Q17 & Q18) Category Effects & Crosstab Analysis

- *Academic Note:* The grouped tables and crosstab functions for Status vs. Location initially returned empty data frames due to the high percentage of missing values in the raw status columns (as identified in Q4). This directly validates the critical need for the advanced data cleaning and predictive modeling pipelines we implemented in the core chapters of this report to reconstruct and predict these missing statuses.

E) Final Reporting

Q20) Final EDA Summary

Top 3 Insights:

1. Transportation distance is a highly variable and significant factor affecting the logistics network.
2. 'Delayed' shipments occur uniformly across various distance ranges, disproving the assumption that only long routes get delayed.
3. The raw data contains significant structural gaps, emphasizing the need for robust data engineering.

Top 3 Problems/Risks:

1. Dates stored as strings and missing actual delivery dates for failed/stuck shipments.
2. Extreme distance outliers could skew simple regression models if not handled properly.
3. High missing value rates in the target variable (ontime/delay).

Next Steps (Implemented in Chapter 5):

- **Cleaning:** Standardize text features and impute missing values.
- **Feature Engineering:** Extract accurate Transit_Hours from raw date columns.
- **Modeling:** Build a Machine Learning classifier (Random Forest) in Python to predict shipment delays accurately.